Music Information Retrieval

http://www.ifs.tuwien.ac.at/mir

Alexander Schindler
Institute of Software Technology and Interactive Systems
Vienna University of Technology
schindler@ifs.tuwien.ac.at

AUSTRIAN INSTITUTE OF TECHNOLOGY
TOMORROW TODAY

FACULTY OF !NFORMATICS

What is Music IR?

- Searching for Music
  - Searching for music on the Web
  - Query by Humming
  - Similarity Retrieval
  - Identity detection (fingerprinting)

- Extraction of information from music

  → plenty of other tasks!

What is Music?

- Music

  Audio: wav, au, mp3, ...

  Symbolic: MIDI, mod, ...

  Scores: Scan, MusicXML

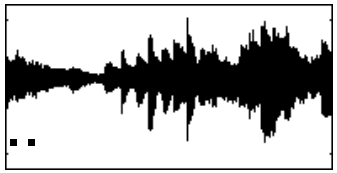

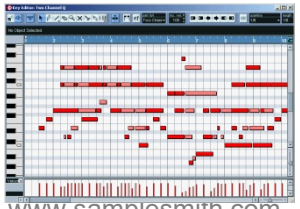

www.samplesmith.com

www.westminster.gov.uk

- Text
  - Song lyrics
  - Artist Biographies
  - Websites: Fanpages, Blogs, Album Reviews, Genre descriptions
- Community data
  - Market basket
  - Tags
  - Social Networks
    - Spotify
    - Last.fm
- Video/Images
  - Album covers
  - Music videos

2.
Feature Extraction from Music

Too much Audio Data

- Digital Audio
  - Sampling Rate: 44,100 Hz
  - 16-bit resolution for each channel
    - 2 channels for stereo
  - 88,200 Integers per second

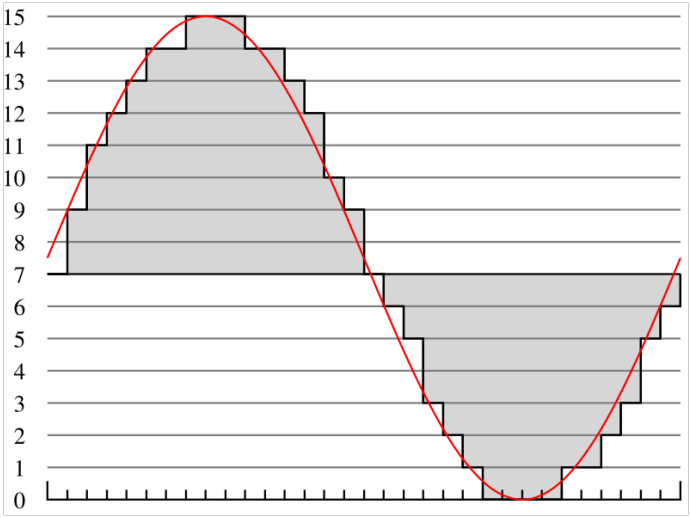

Excercise: Find Documents Containing the Word „Music“

Document 1:

*“Most of these issues stem from the commercial interest in **music** by record labels, and therefore imposed rigid copyright issues, that prevent researchers from sharing their **music** collections with others. Subsequently, only a limited number of data sets has risen to a pseudo benchmark level, i.e. where most of the researchers in the field have access to the same collection.“*

Document 2:

*“The Echonest Analyzer [5] is a **music** audio analysis tool available as a free Web service accessible over the Echonest API and as a commercially distributed standalone command line tool. The Analyzer implements an onset detector which is used for segmentation.”*

Document 3:

*“The Million Song Dataset (MSD), a collection of one million **music** pieces, enables a new era of research of **Music** Information Retrieval methods for large-scale applications. It comes as a collection of meta-data such as the song names, artists and albums, together with a set of features extracted with the The Echo Nest services, such as loudness, tempo, and MFCC-like features.”*

Excercise: Find Songs with Strings

Song 1:

83, 58, 11, 11, 9, 60, 96, 25, 39, 42, 87, 90, 12, 26, 99, 69, 10, 56, 64, 41, 47, 61, 6, 40, 94, 23, 43, 52, 31, 77, 32, 57, 40, 89, 91, 28, 38, 96, 3, 90, 43, 18, 25, 16, 79, 97, 83, 64, 46, 70, 63, 34, 38, 39, 7, 66, 89, 95, 9, 47, 11, 59, 9, 17, 46, 92, 27, 58, 87, 46, 39, 100, 10, 2, 5, 53, 73, 56, 43, 46, 47, 67, 2, 60, 9, 23, 43, 21, 98, 34, 29, 62, 26, 72, 38, 98

Song 2:

55, 96, 11, 49, 83, 58, 11, 11, 9, 60, 96, 25, 39, 42, 87, 90, 12, 26, 99, 69, 10, 56, 64, 41, 47, 61, 6, 40, 94, 23, 43, 52, 31, 77, 32, 57, 40, 89, 91, 28, 38, 96, 3, 90, 43, 18, 25, 16, 79, 97, 83, 64, 46, 70, 63, 34, 38, 39, 7, 66, 89, 95, 9, 47, 11, 59, 9, 17, 46, 92, 27, 58, 87, 46, 39, 100, 10, 2, 5, 53, 73, 56, 43, 46, 47, 67, 2, 60, 9, 23, 43, 21, 98, 34, 29, 62, 26, 72, 38, 98, 55, 96, 11, 49, 83, 58, 11, 11, 9, 60, 96, 25, 39, 42, 87, 90, 12, 26, 99, 69, 10, 56, 64, 41, 47, 61, 6, 40, 94, 23, 43, 52, 31, 77, 32, 57, 40, 89, 91, 28, 38, 96, 3, 90, 43, 18, 25, 16, 79, 97, 83, 64, 46, 70, 63, 34, 38, 39, 7

Song 3:

66, 89, 95, 9, 47, 11, 59, 9, 17, 46, 92, 27, 58, 87, 46, 39, 100, 10, 2, 5, 53, 73, 56, 43, 46, 47, 67, 2, 60, 9, 23, 43, 21, 98, 34, 29, 62, 26, 72, 38, 98, 55, 96, 11, 49, 83, 58, 11, 11, 9, 60, 96, 25, 39, 42, 87, 90, 12, 26, 99, 69, 10, 56, 64, 41, 47, 61, 6, 40, 94, 23, 43, 52, 31, 77, 32, 57, 40, 89, 91, 28, 38, 96, 3, 90, 43, 18, 25, 16, 79, 97, 83, 64, 46, 70, 63, 34, 38, 39, 7, 66, 89, 95, 9, 47, 11, 59, 9, 17, 46, 92, 27, 58, 87, 46, 39, 100, 10, 2, 5, 53, 73, 56, 43, 46, 47, 67, 2

Excercise: Same Genre?

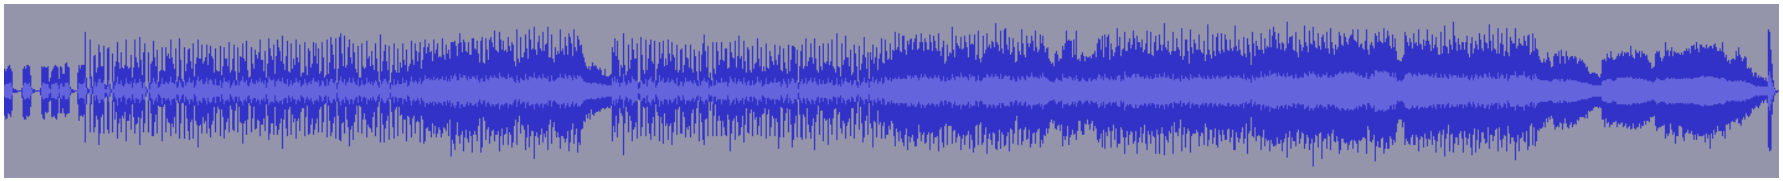

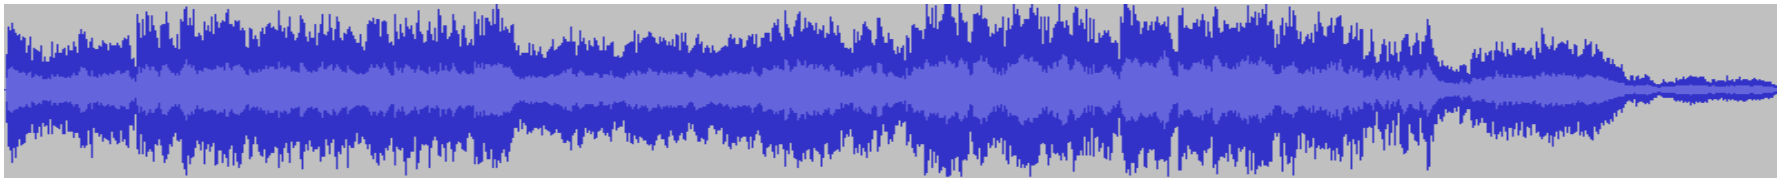

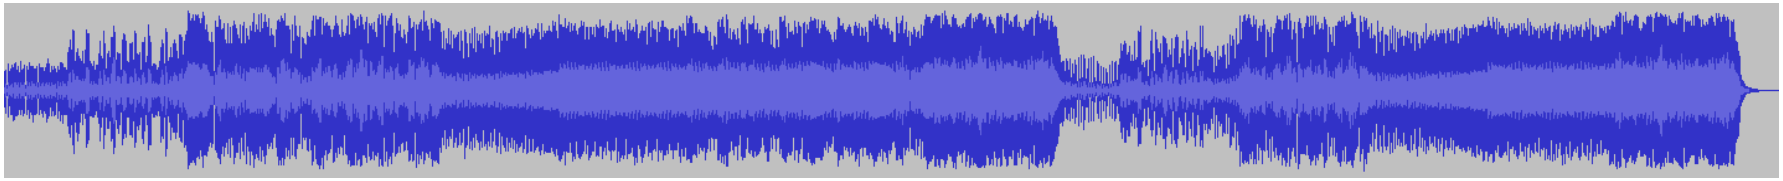

Excercise: Identify Songs

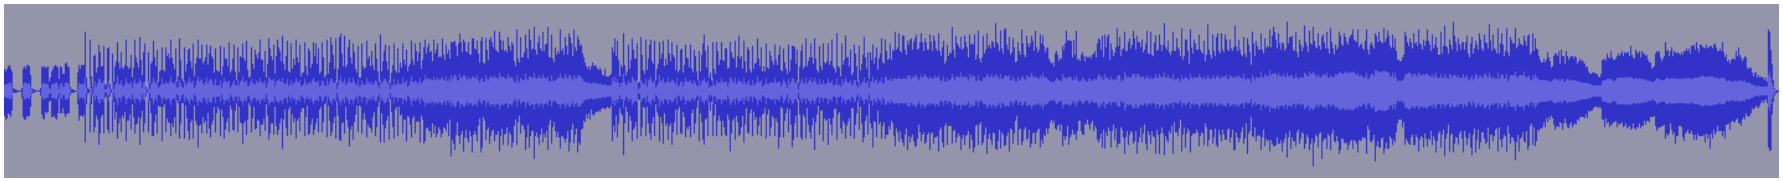

AC-DC – Highway to Hell

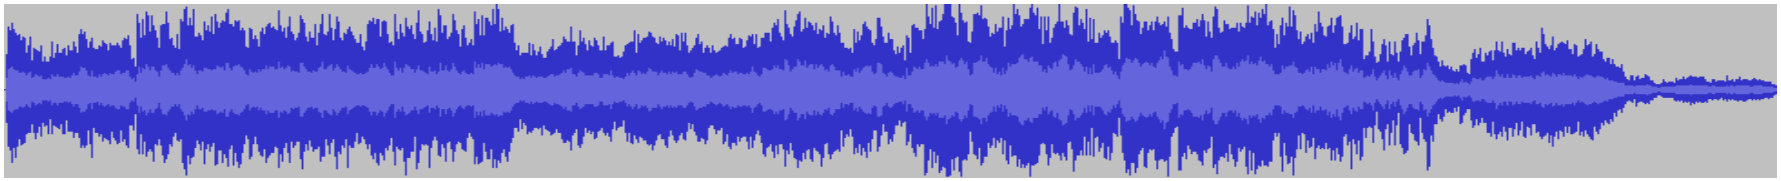

John Williams – Star Wars Main Theme

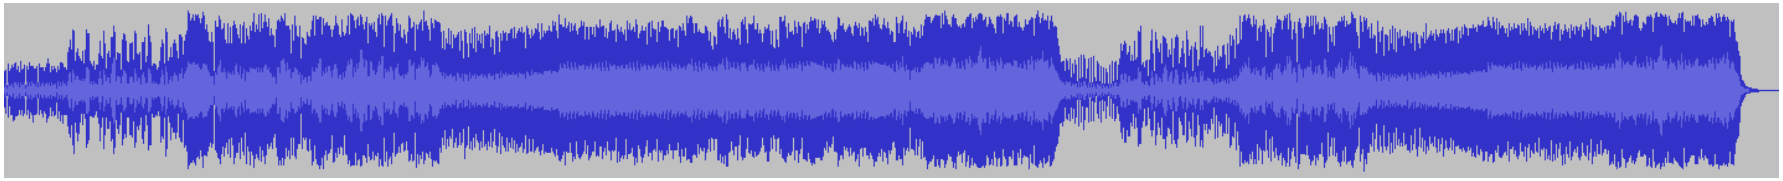

Rihanna feat. Calvin Harris – We Found Love

Audio Feature Extraction

- Reduce audio data by extracting information about:
  - Pitch
  - Timbre
  - Rhythm
  - etc.
- → extract „audio descriptors“

Problem: Source Separation

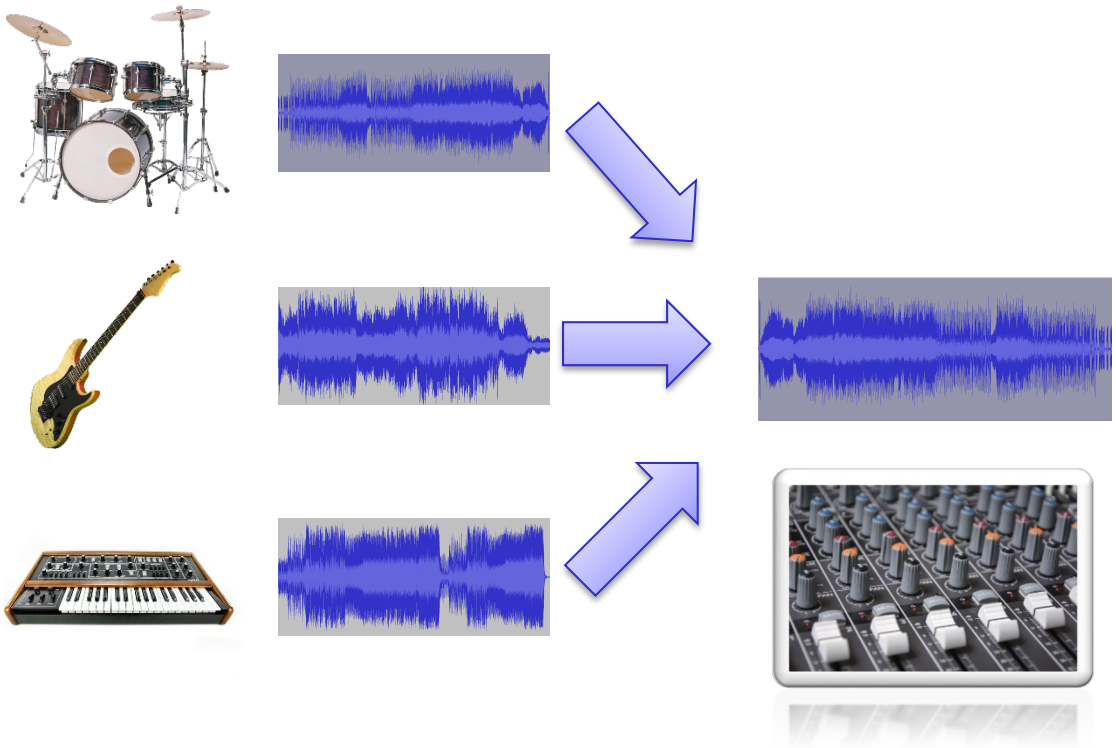

Audio Feature Extraction

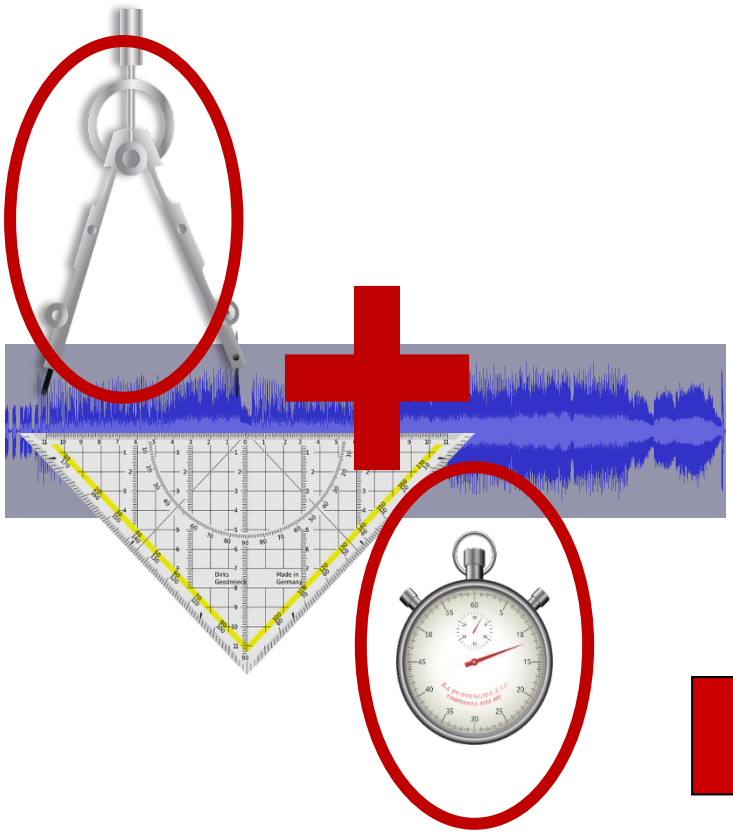

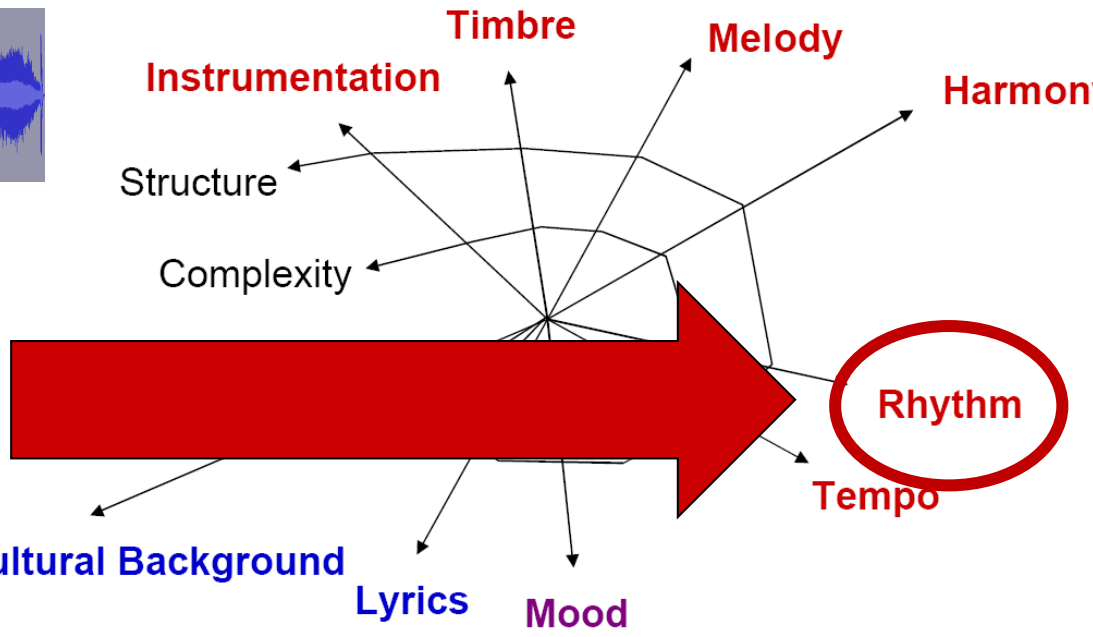

Signal Processing

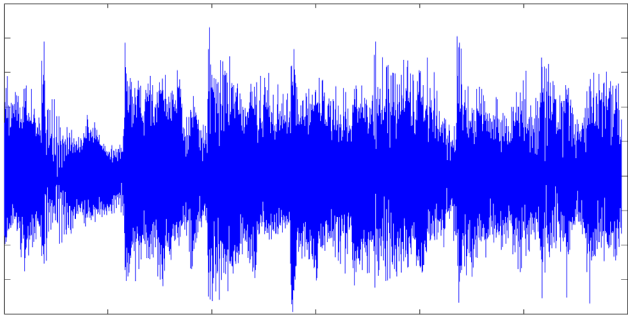

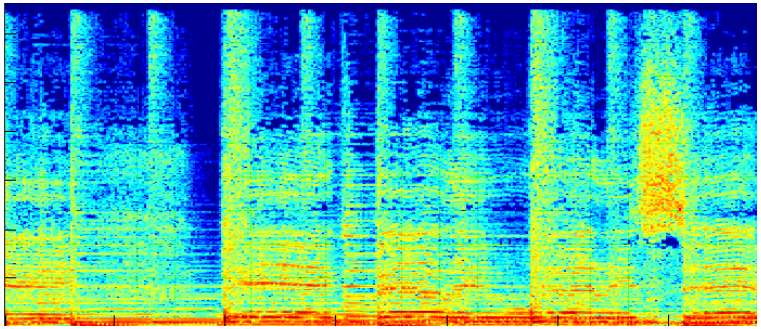

Time Domain
(„Wave Form“)

Frequency Domain
(„Spectrum“)

Time-Frequency Transformation

Fourier Transform (FFT)
Discrete Cosine Transform (DCT)
Wavelet Transform

Fourier Transform

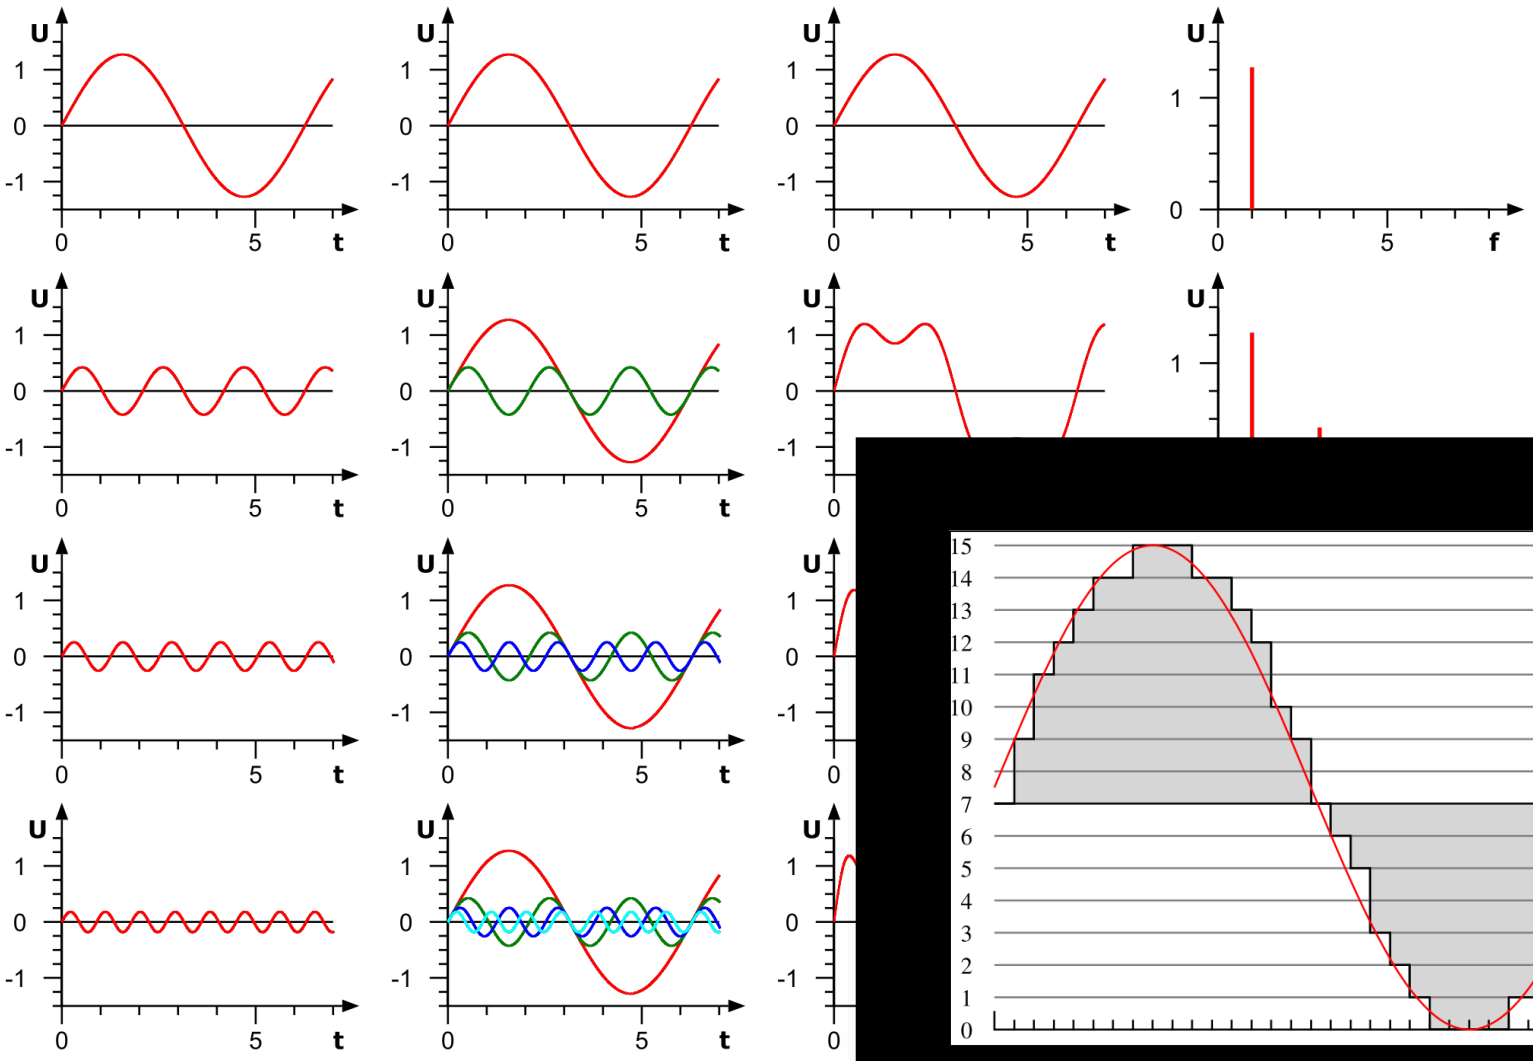

Audio/Music Feature Extraction
by example…

Rhythm Pattern (RP)

- fluctuations on critical frequency bands
  (a.k.a. Fluctuation Pattern)

- covers rhythm in the broad sense

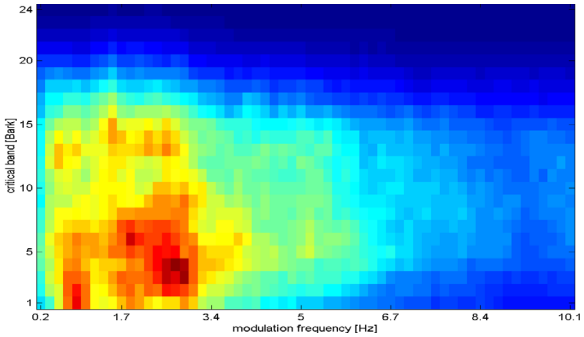

Classical

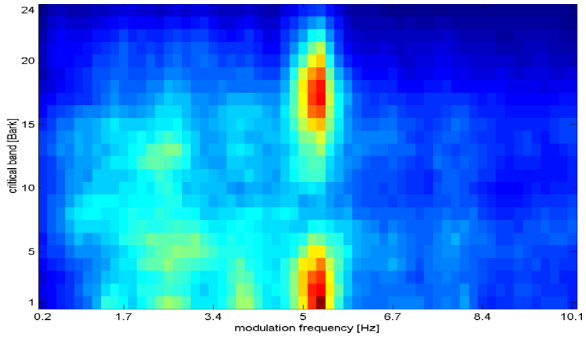

Rock

Rhythm Pattern (RP)

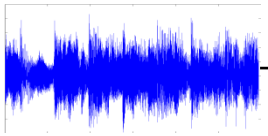

Segmentation (6 s)

Spectrogram (FFT)

Bark scale (24 bands)

Spectral Masking

Decibel

Phon

Sone

Modulation Amplitudes

Fluctuation Strength

Filtering/Smoothing

Rhythm Pattern features

(1440 dimensions)

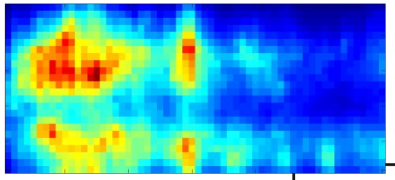

Bark Scale

- psychoacoustical scale (related to Mel scale)
- 24 „critical bands“ of hearing (non-linear)
- proposed by Eberhard Zwicker in 1961

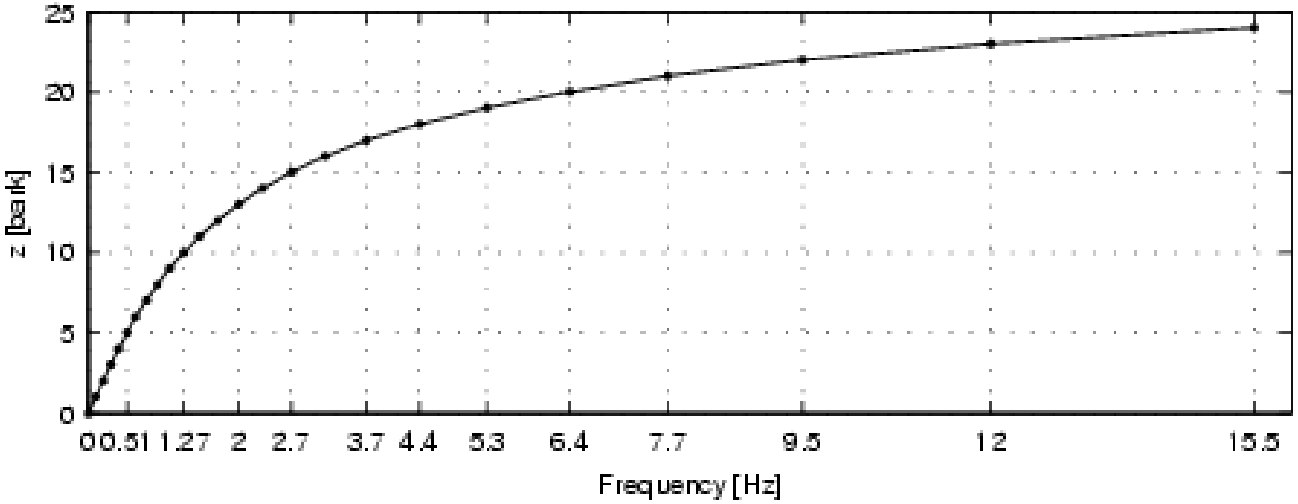

http://en.wikipedia.org/wiki/Bark_scale

Equal loudness curves (Phon)

- Relationship between sound pressure level in decibel and hearing sensation is not linear
- Perceived loudness depends on frequency of the tone
- equal loudness contours for 3, 20, 40, 60, 80, 100 phon

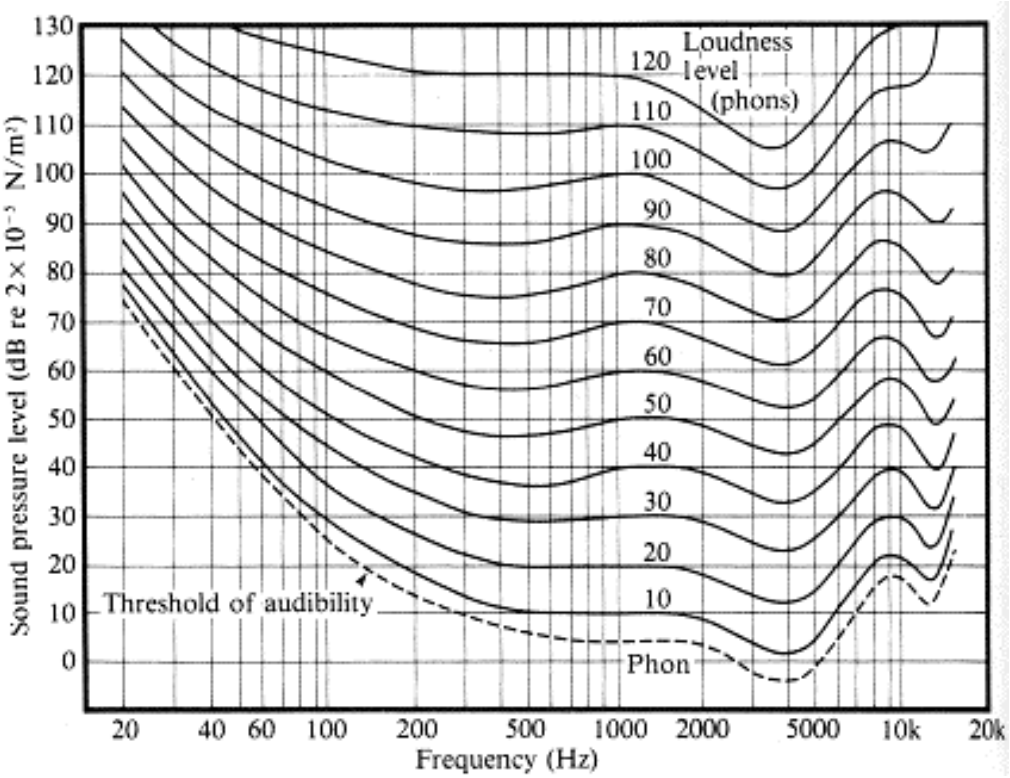

on-line test: http://www.phys.unsw.edu.au/jw/hearing.html

Sone Transformation

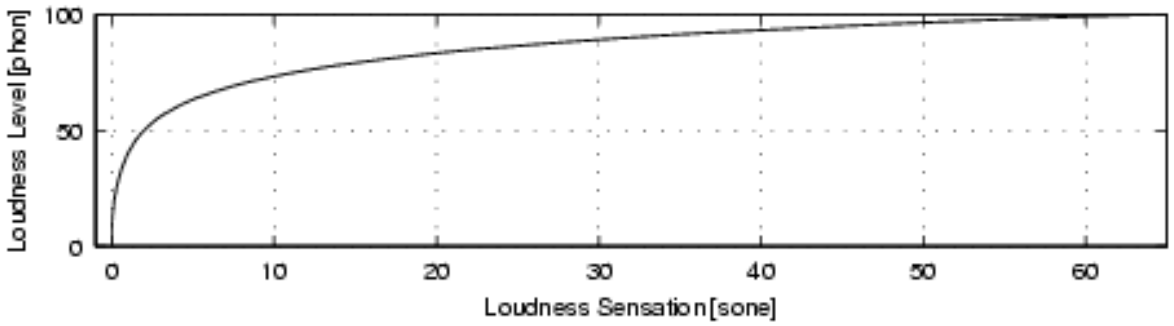

Sone	1	2	4	8	16	32	64
Phon	40	50	60	70	80	90	100

- Perceived loudness measured in Phon does not increase linearly
- Transformation into Sone
- Up to 40 phon slow increase in perceived loudness, then drastic increase
- Higher sensibility for certain loudness differences

Rhythm Pattern (RP): 2 examples

Queen – Another One Bites The Dust (first 6 seconds)

PCM Audio Signal

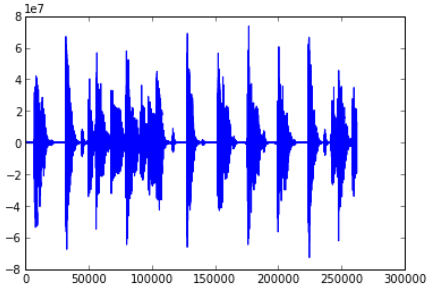

Power Spectrum

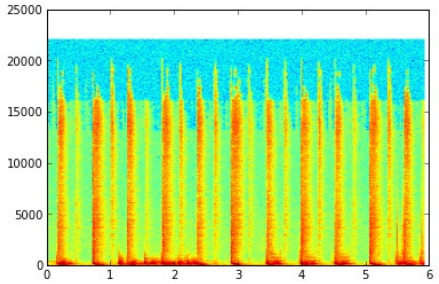

Bark Scale

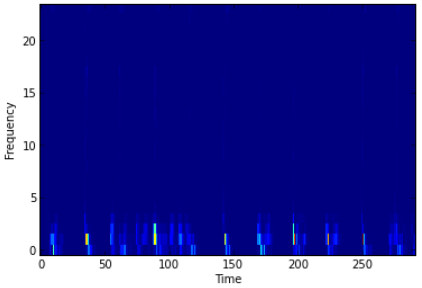

Decibel

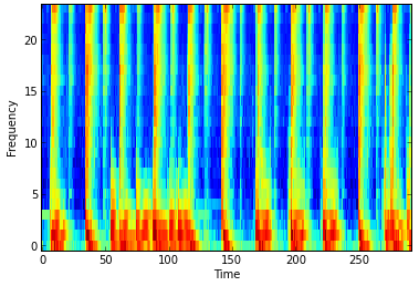

Phon

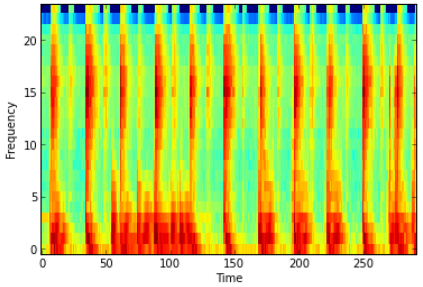

Sone

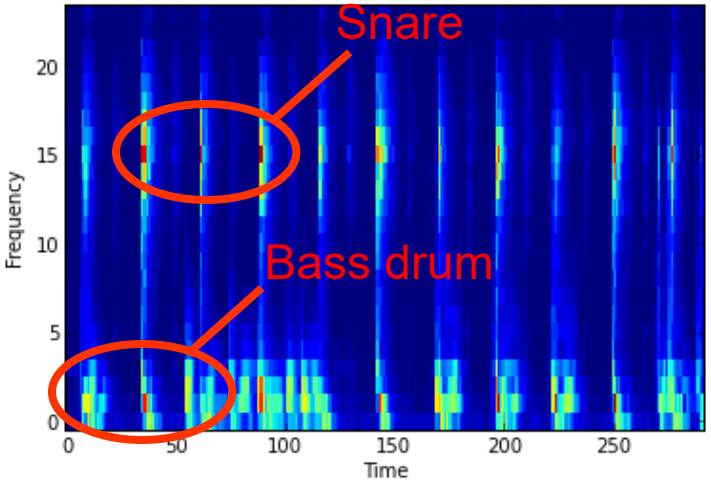

Rhythm Pattern (RP): 2 examples

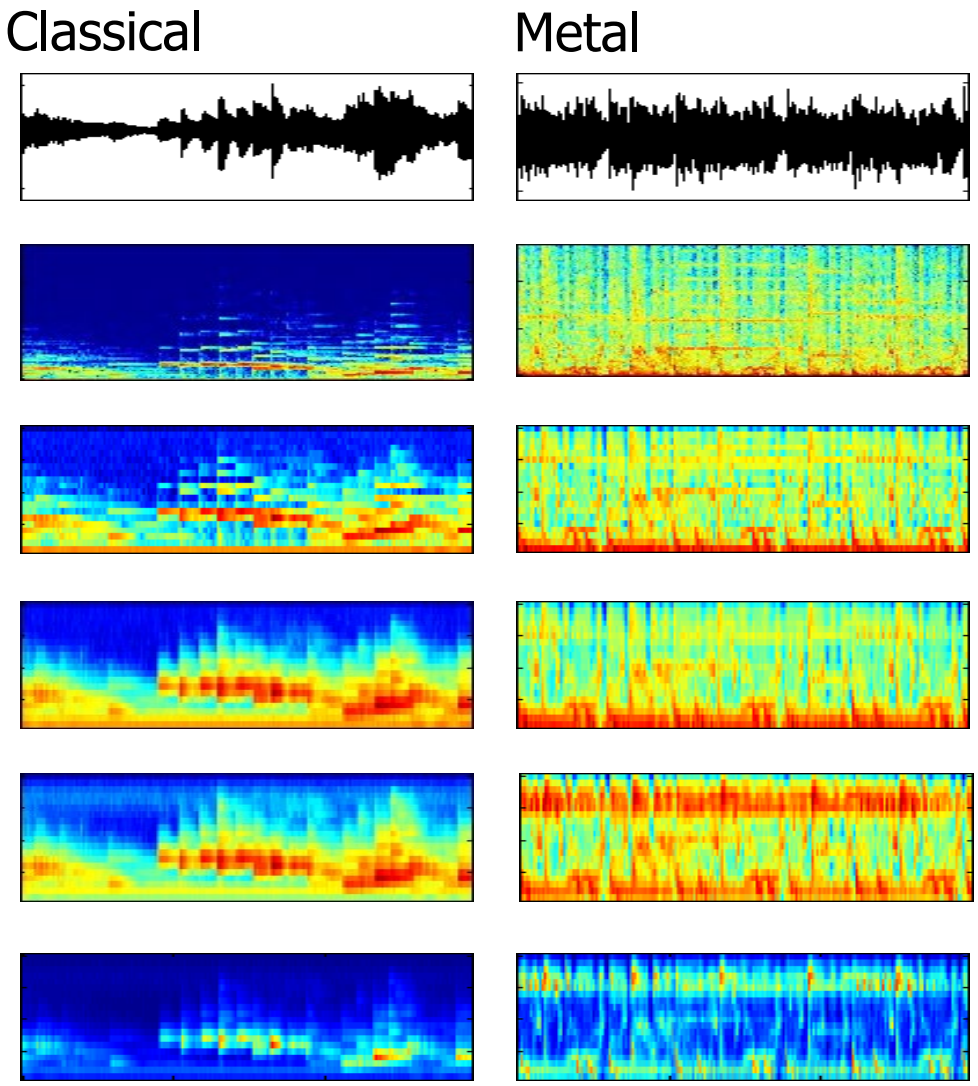

Rhythm Pattern (RP): 2 examples

Classical

Metal

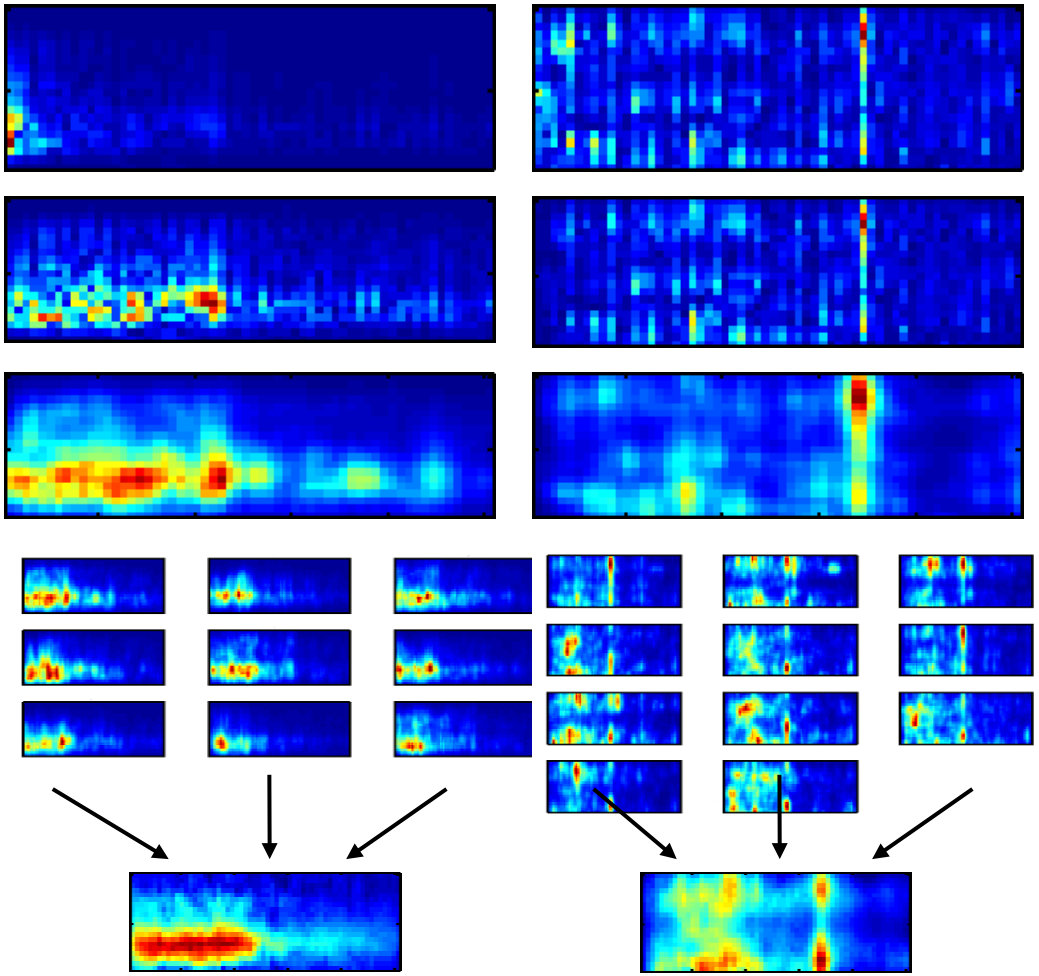

modulation amplitude spectrum (“cepstrum”)

Fluctuation Strength

Filter (Gradient, Gauss)

Median

24*60=
1.440-dim feature vec.

RP per Genre

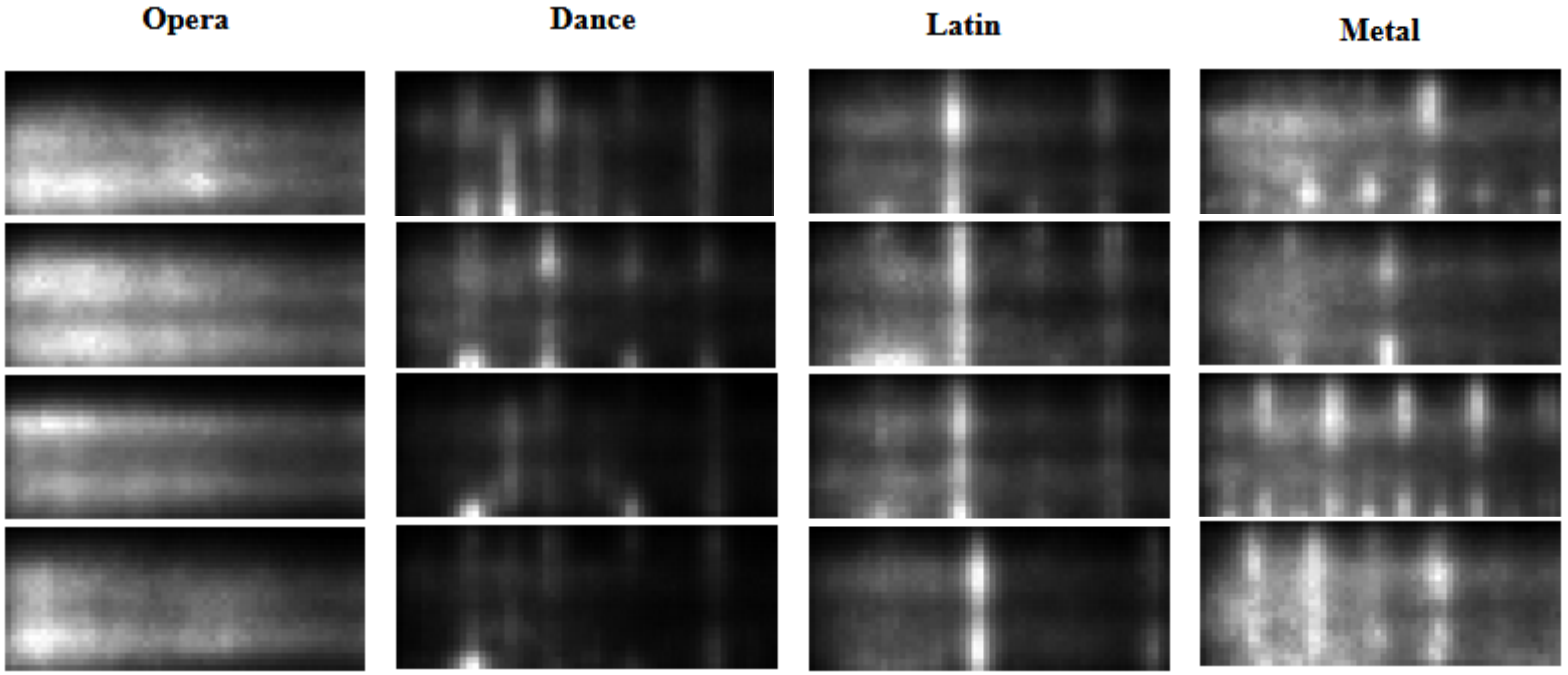

Modulated Synthesizer

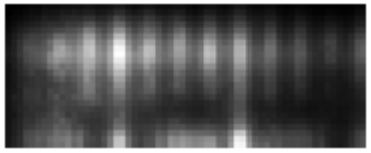

Statistical Spectrum Descriptors

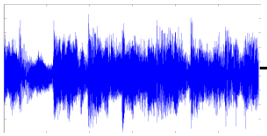

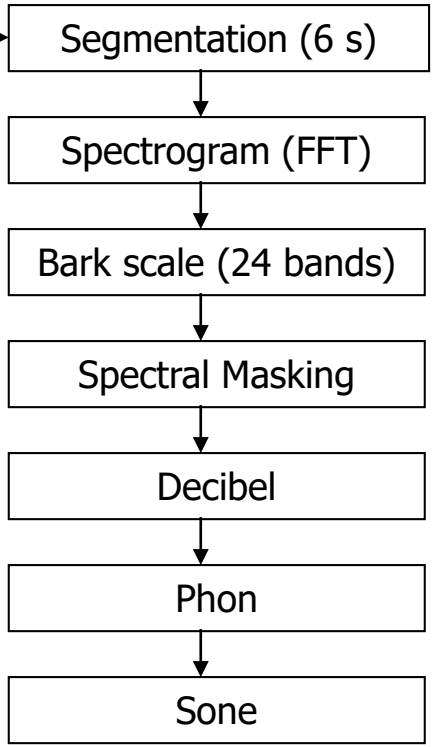

- description of each of the 24 critical bands of the Sonogram by 7 statistical measures
- 168 feature attributes (24x7)

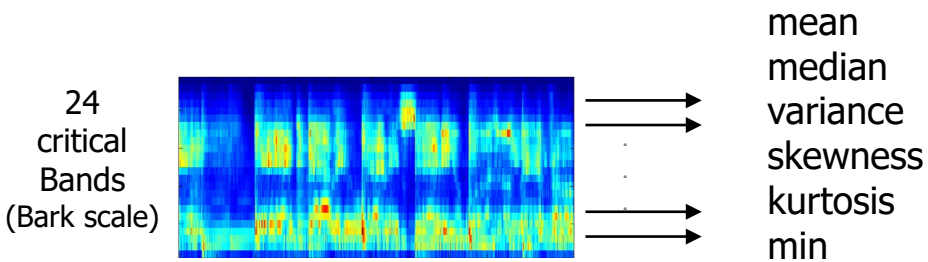

Bark-scale Sonogram (after Sone Step of RP)

Deep Learning
for
Music Information Retrieval

Convolutional Neural Network (CNN)

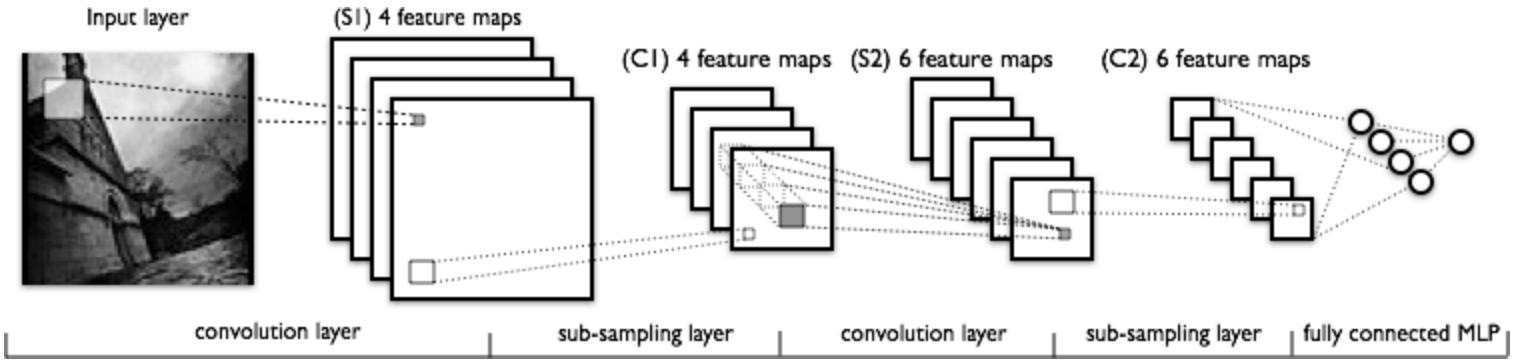

http://deeplearning.net/tutorial/lenet.html

Combines three types of layers:

- **Convolutional layer:** performs 2D convolution of 2D input with multiple learned 2D kernels
- **Subsampling layer:** replaces 2D patches by their maximum ("max-pooling") or average
- **Fully-connected layer:** computes weighted sums of its input with multiple sets of learned coefficients

Applies a nonlinear function after each linear operation (without, a deep network would be linear despite its depth).

Full vs. Convolutional Layer / Network

Fully-connected layer:
Each **input** is a **scalar** value,
each **weight** is a **scalar** value,
each output is the sum of
inputs **multiplied** by weights.

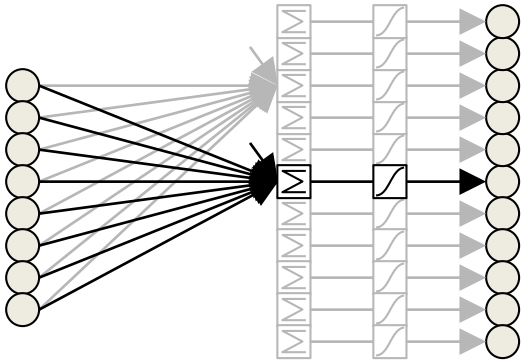

Convolutional layer:
Each **input** is a **tensor** (e.g., 2D),
each **weight** is a **tensor**,
each output is the sum of
inputs **convolved** by weights.

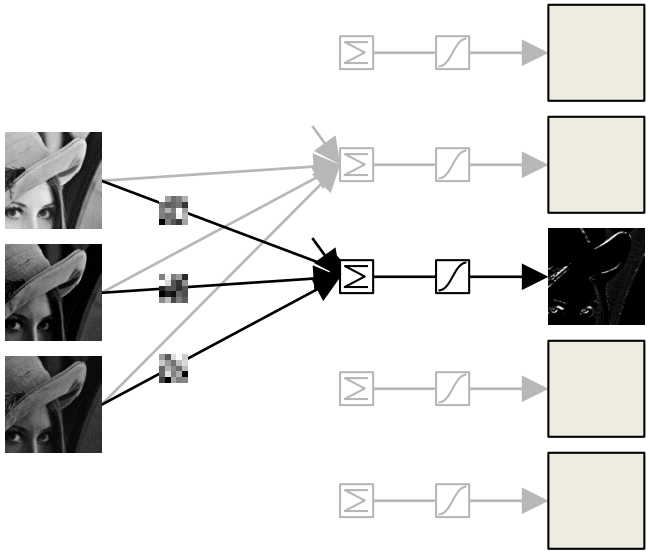

Motivation for Convolutions

- Apply local filter kernels
- These kernels are the neurons that are learned

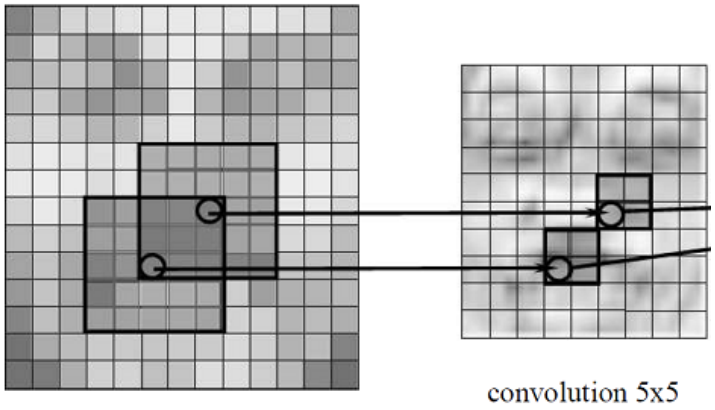

Operation	Kernel	Image result
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	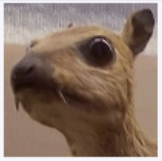
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	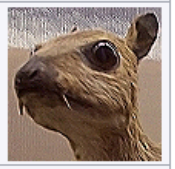
Box blur (normalized)	$\frac{1}{9}\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	

Images: http://sanghyukchun.github.io/75/
https://en.wikipedia.org/wiki/Kernel_(image_processing)

Deep Learning for Music IR

Pre-Processing: Waveform → Spectrogram → 40 Mel bands → Log scale

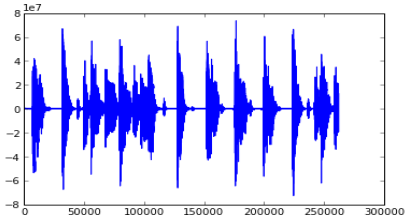

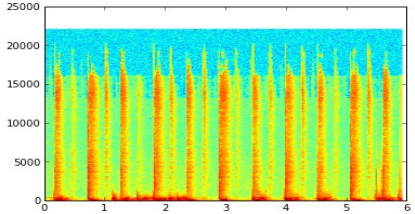

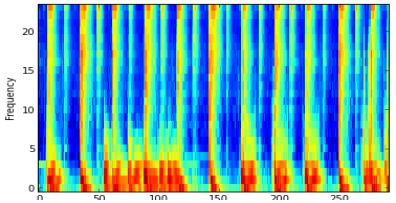

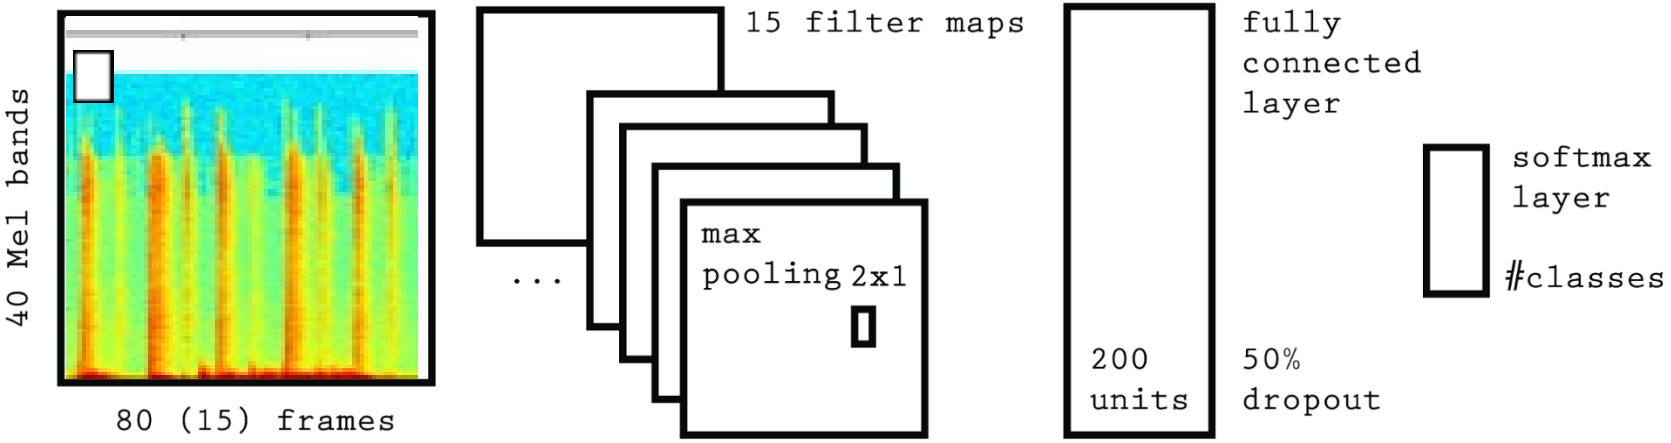

Winning algorithm MIREX 2015 music/speech classification task (99.73%) by Thomas Lidy

Visualizing CNN Filters

learned for Music/Speech Classification

Learned Filter Weights

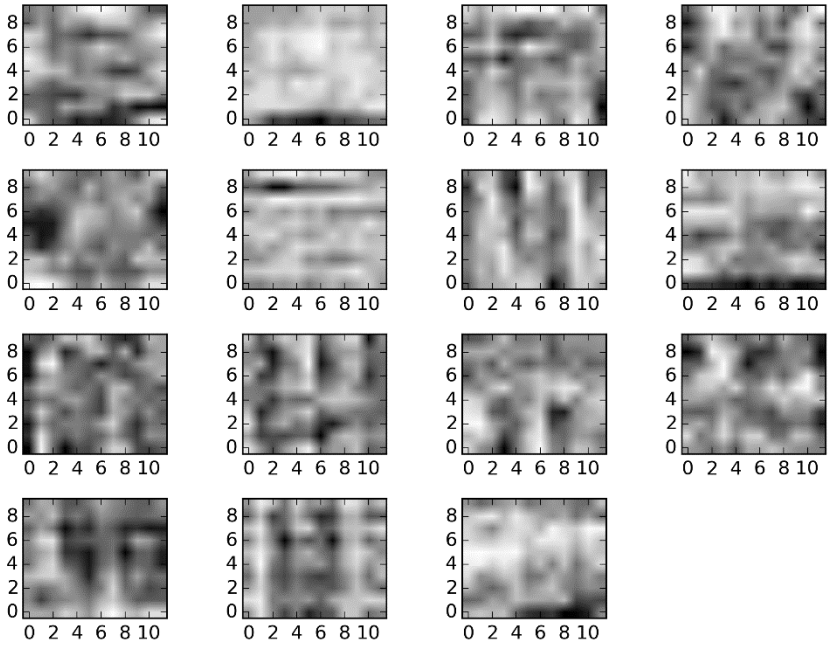

Convolved Spectrograms

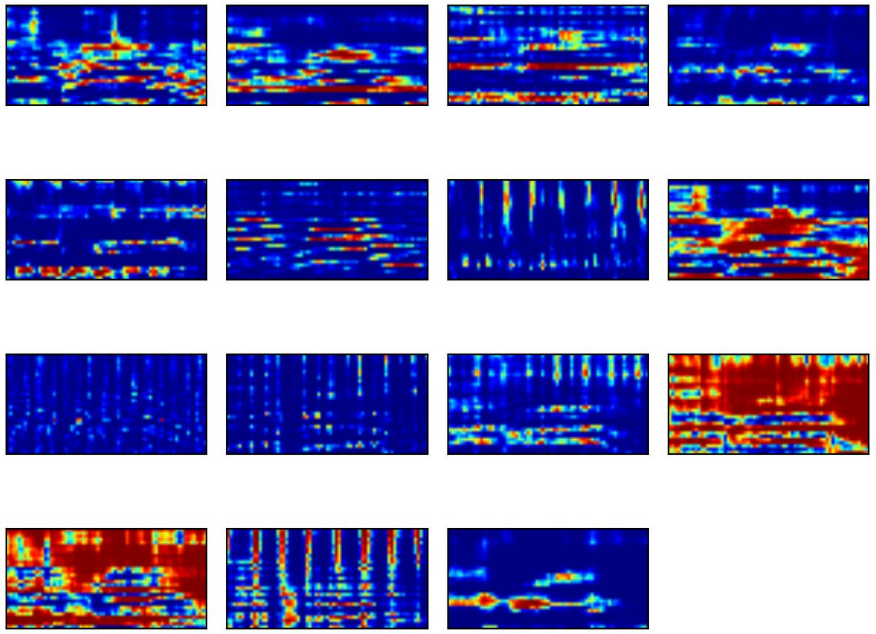

Audio and Music related Research at IFS

Similarity Search

Search for similar sounding Audio Content in Europeana

- Select a favoured track
- System analyzes audio content
- Provides a list of tracks with calculated similar accustic properties

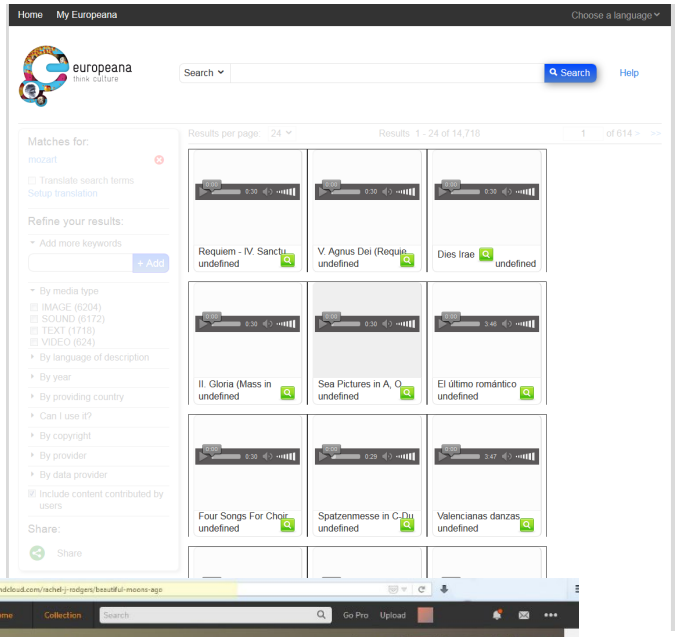

Search for similar sounding Soundcloud tracks in Europeana

- Supply a Soundcloud URL
- System downloads and analyzes track
- Provides a list of similar sounding Europeana tracks

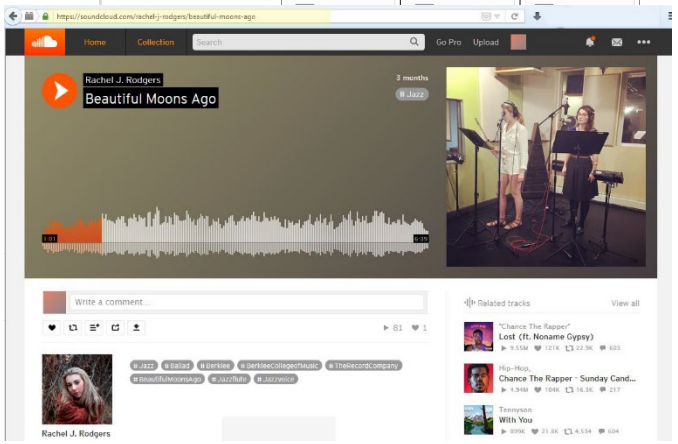

Score Following

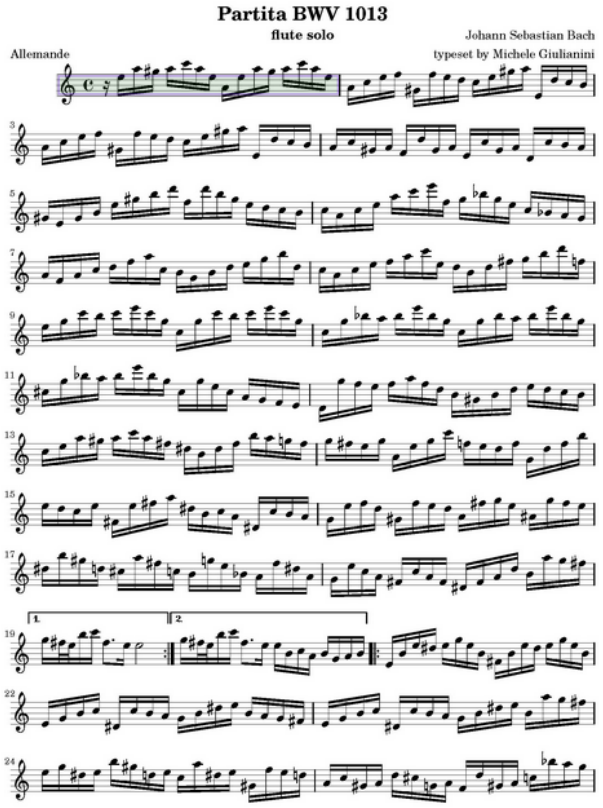

http://www.ifs.tuwien.ac.at/~schindler/files/eusounds/scorefollowing/SFP_Bach_BWV_1013.html

MVIR Objectives

- Multimodal Approach to MIR Problems
  - Classification / Tagging
  - Mood estimation
  - Music Similarity Retrieval

- Hypothesis:
  - visual layer of music videos conta
    related information

- Research Aims
  - Can this information be used?
    - Improve MIR solutions
    - Use images as queries

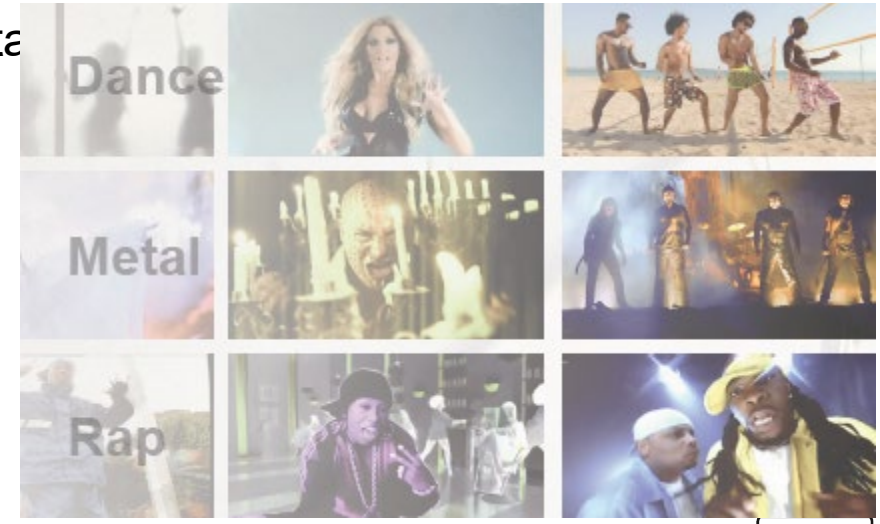

Top concepts of music video frames examples

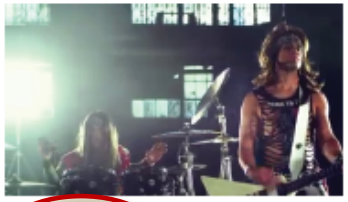

stage	0.3162
electric guitar	0.1169
bassoon	0.0649
accordion	0.0611
drumstick	0.0386
microphone	0.0313
marimba	0.0276

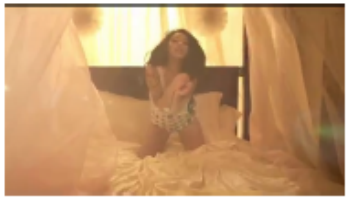

mosquito net	0.0932
wardrobe	0.0857
brassiere	0.0815
shower curtain	0.0471
candle	0.0400
plastic bag	0.0204
hoopskirt	0.0187

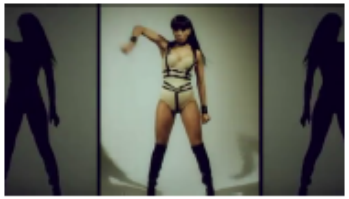

maillot	0.2745
bolo tie	0.0732
Windsor tie	0.0550
letter opener	0.0486
brassiere	0.0390
bikini	0.0384
bassoon	0.0364

lumbermill	0.1925
tow truck	0.1215
harvester	0.1152
thresher	0.0513
jeep	0.0484
half track	0.0473
pickup truck	0.0460

wig	0.4399
neck brace	0.0577
chimpanzee	0.0418
hair spray	0.0375
orangutan	0.0366
cloak	0.0267
Windsor tie	0.0236

Classification results (visual concepts only)

(c) High-level Visual Concepts											
$v_{in}1$	MEAN	1000	66.86	42.09	53.69	51.26	31.23	37.05	46.87	23.90	**33.07**
$v_{in}2$	STD	1000	69.78	46.76	50.08	51.95	29.99	32.88	48.29	26.83	29.63
$v_{in}3$	MAX	1000	73.15	44.26	46.41	54.60	33.05	31.94	50.07	26.93	27.49
$v_{in}4$	$v_{in}3$+$v_{in}2$	2000	73.61	46.53	**51.21**	55.04	31.48	34.00	51.30	27.03	31.04
$v_{in}5$	$v_{in}3$+$v_{in}1$	2000	**74.36**	47.70	53.65	55.99	**33.70**	**37.83**	**51.58**	**28.88**	33.83

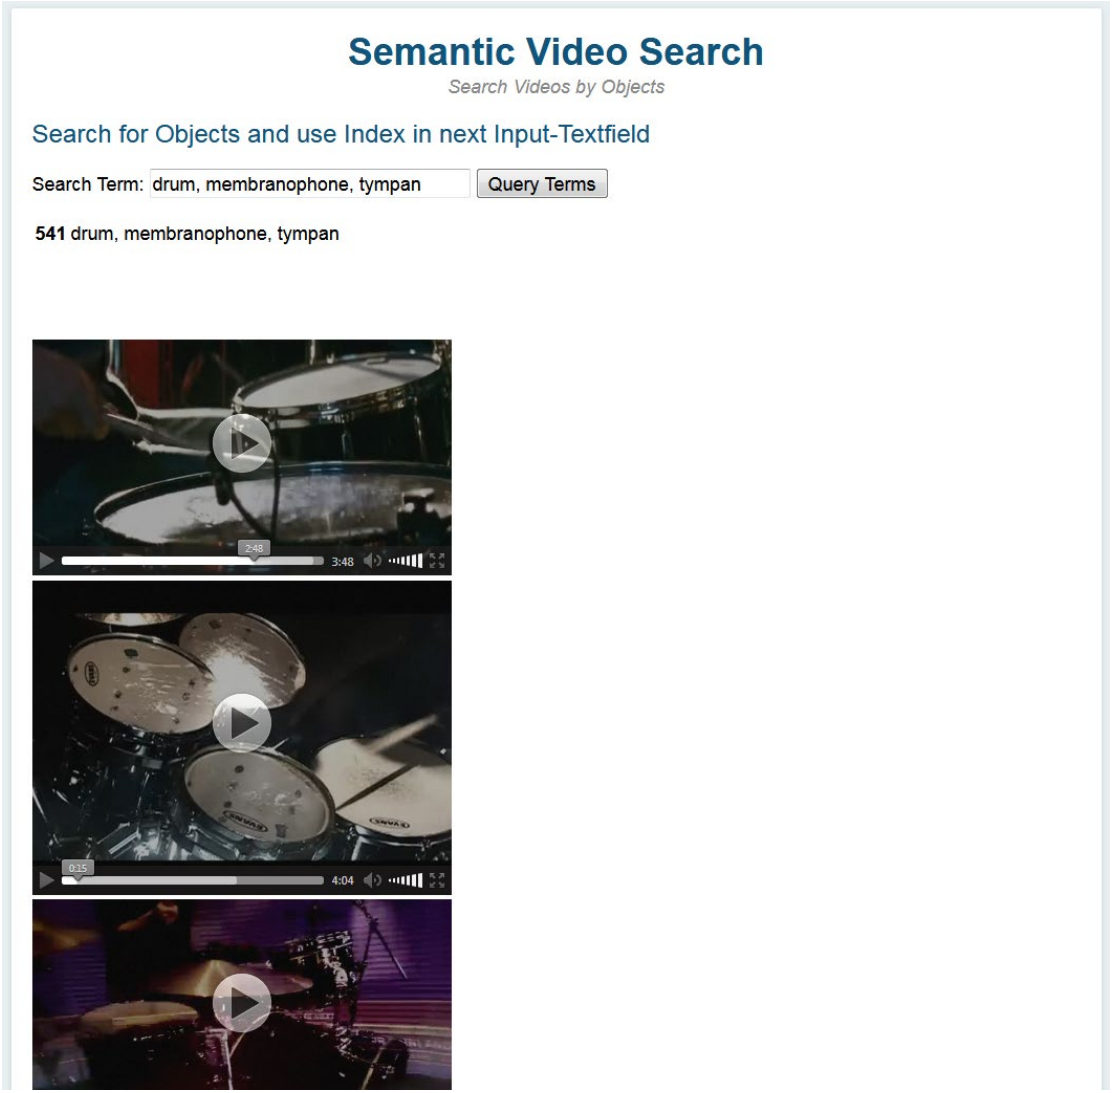

http://172.20.36.10:5000/

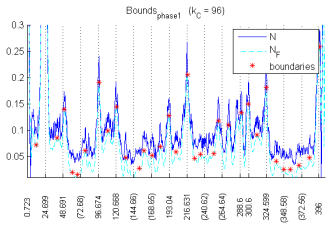

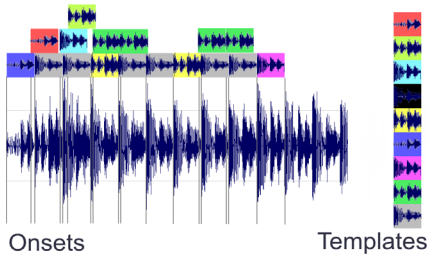

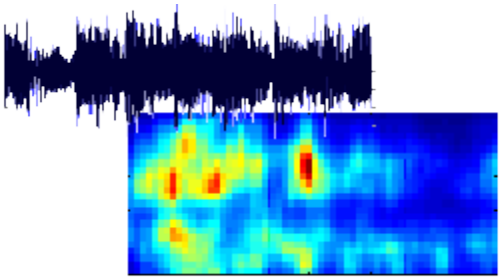

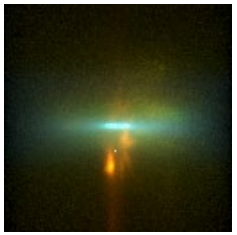

Thank You !

Alexander Schindler - schindler@ifs.tuwien.ac.at

http://www.ifs.tuwien.ac.at/mir