LARGE SCALE AUDIO-VISUAL VIDEO ANALYTICS PLATFORM

FOR

FORENSIC INVESTIGATIONS

OF

TERRORISTIC ATTACKS

Alexander Schindler
Martin Boyer
Andrew Lindley
David Schreiber
Thomas Philipp

Alexander Schindler
Scientist
Information Management
Center For Digital Safety & Security

AIT Austrian Institute Of Technology Gmbh
Alexander.schindler@ait.ac.at

LARGE SCALE AUDIO-VISUAL VIDEO ANALYTICS PLATFORM

FOR

FORENSIC INVESTIGATIONS

OF

TERRORISTIC ATTACKS

Alexander Schindler
Martin Boyer
Andrew Lindley
David Schreiber
Thomas Philipp

Alexander Schindler
Scientist
Information Management
Center For Digital Safety & Security

AIT Austrian Institute Of Technology Gmbh
Alexander.schindler@ait.ac.at

FORENSIC INVESTIGATIONS OF TERRORISTIC ATTACKS

AUSTRIAN INSTITUTE OF TECHNOLOGY
TOMORROW TODAY

- **Context**
  - Forensic Investigation
  - Investigating video data after a terroristic attack

- **Objectives**
  - Spot suspects
  - Follow hints by civilian witnesses
  - Collect and secure evidence
  - Prevent immediate or subsequent attacks

FORENSIC INVESTIGATIONS OF TERRORISTIC ATTACKS

AUSTRIAN INSTITUTE OF TECHNOLOGY
TOMORROW TODAY

- **Obstacles**
  - Great increase in the number of public and private cameras
  - Massively increasing volume of video data to be analysed
    - Boston Marathon Bombing 5.000h
    - Toulouse and Montauban:10.000h (35TB)
  - Time pressure
    - Timely content evaluation of video mass data is of considerable importance

FORENSIC INVESTIGATIONS OF TERRORISTIC ATTACKS

AUSTRIAN INSTITUTE OF TECHNOLOGY
TOMORROW TODAY

- **Initial Situation (before project)**
  - manual viewing/processing of the video material
  - Personnel-intensive: time span from several hundred to several thousand hours

  - **Technical, supporting tools necessary**

- **Projects Goals and Outcomes**
  - 2 Projects
    - FLORIDA (Bi-Lateral funding Austria/Germany) => intial research
    - VICTORIA (H2020) => TRL 6 - 10
  - Large-scale computing platform
  - Analytical modules

AUSTRIAN INSTITUTE OF TECHNOLOGY
TOMORROW TODAY

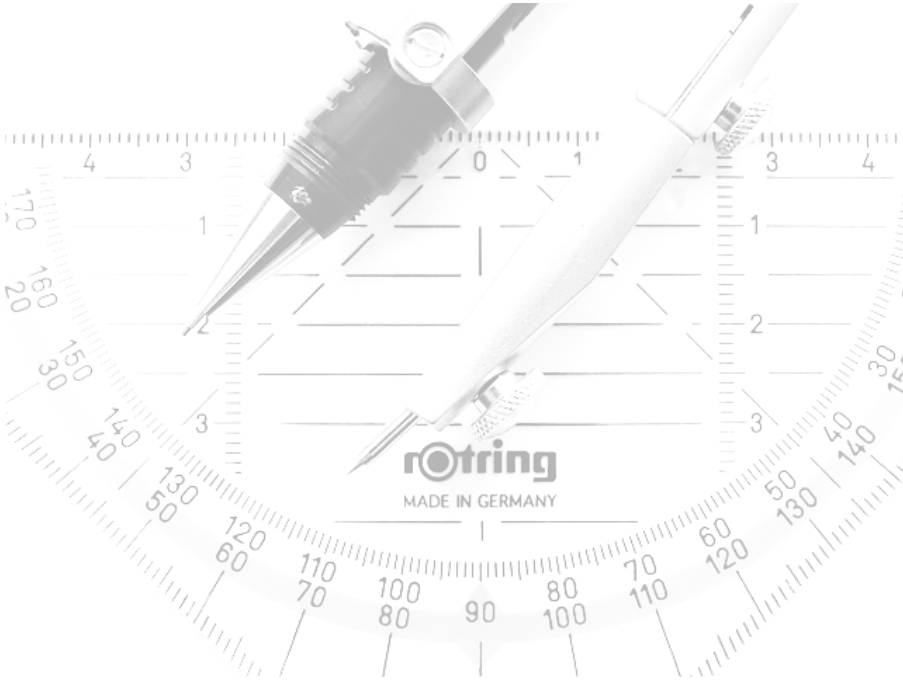

AUDIO EVENT DETECTION

Analytic Modules

AUDIO EVENT DETECTION

- **Task**
  - Detect and predict audio-events into predefined categories
    - Gunshots, explosions, emergency vehicles, scream, speech, Alarm

- **Use-Case**
  - Content filter in mass video-data
  - Example: attack with firearms
    - => initiate search by filtering all videos which contain *Gunshots* (sorted by confidence)

AUDIO EVENT DETECTION

- **Applied Technology**
  - Recurrent Convolutional Neual Networks
  - With Attention Layer

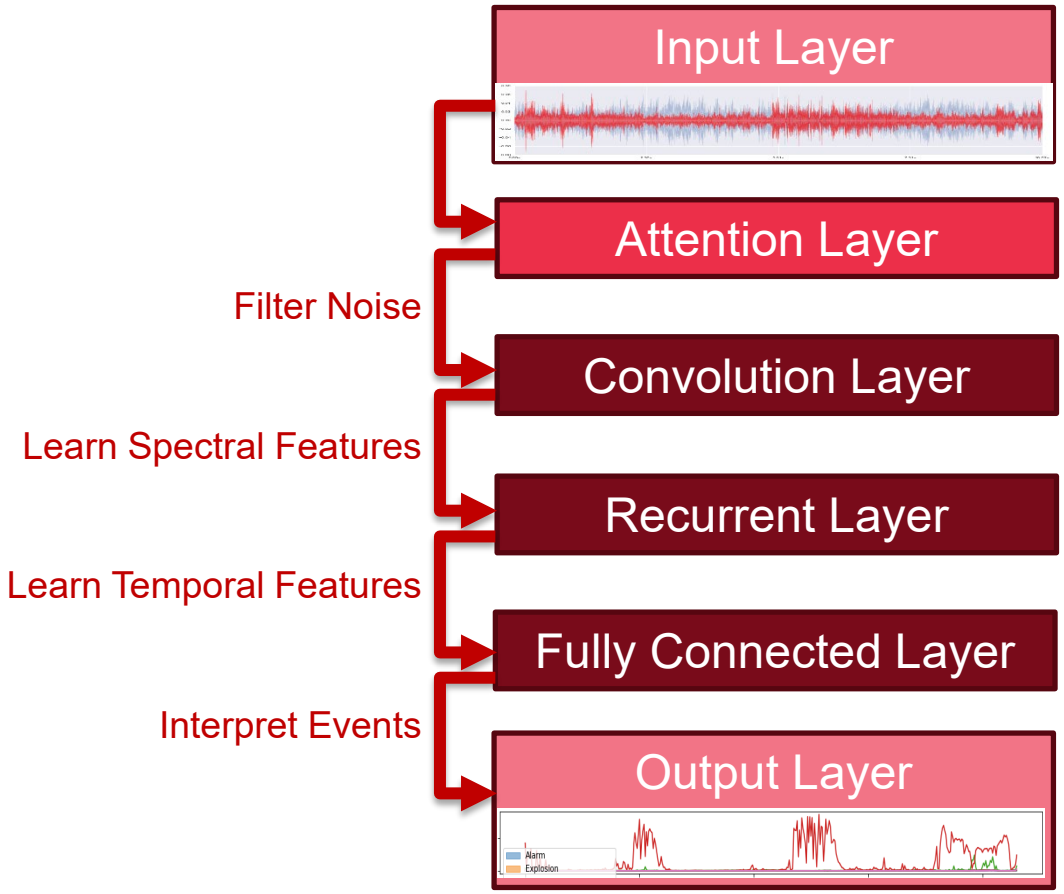

AUDIO EVENT DETECTION

1. **Input representation**
   - **Common: Mel-Spectrograms**

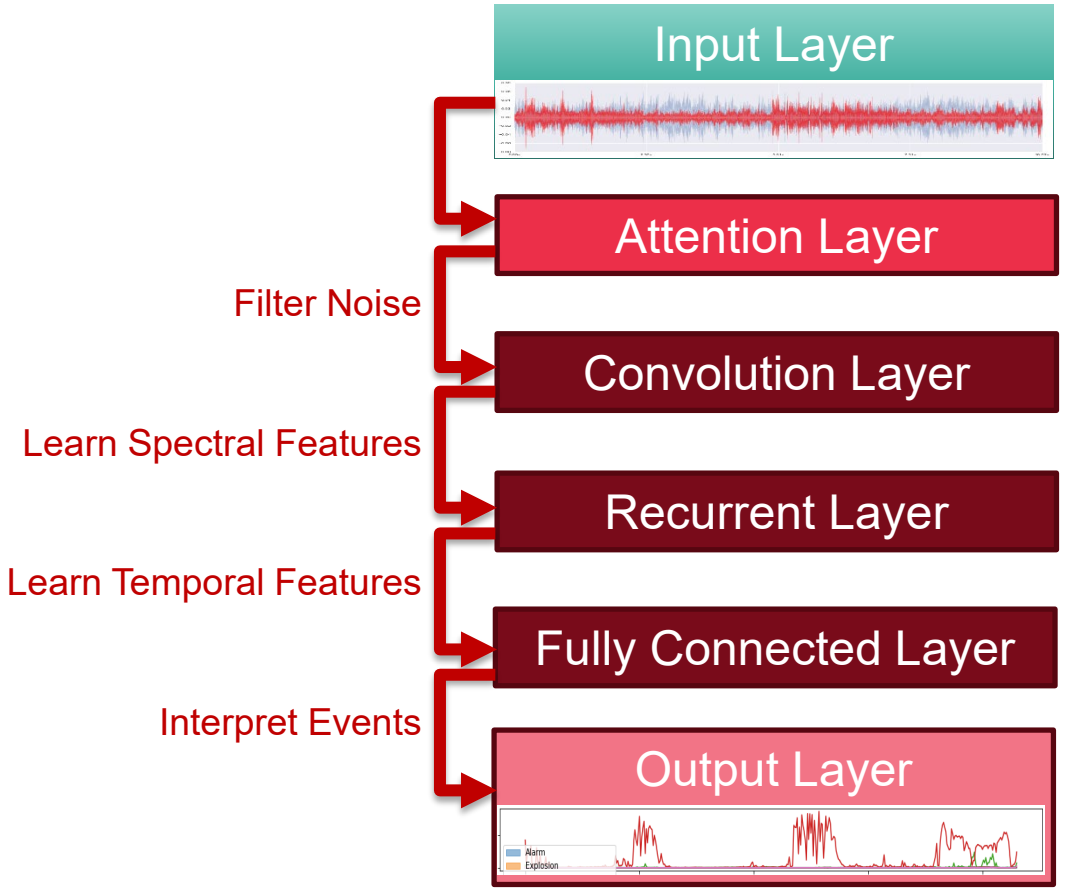

AUDIO EVENT DETECTION

1. Input representation
   - Common: Mel-Spectrograms

- **Attention Layer**
  - **Filter non-relevant information from Input**
  - **Help to learn faster**
  - **Better convergence**
  - **Better generalization**
  - **Smoother prediction signal**

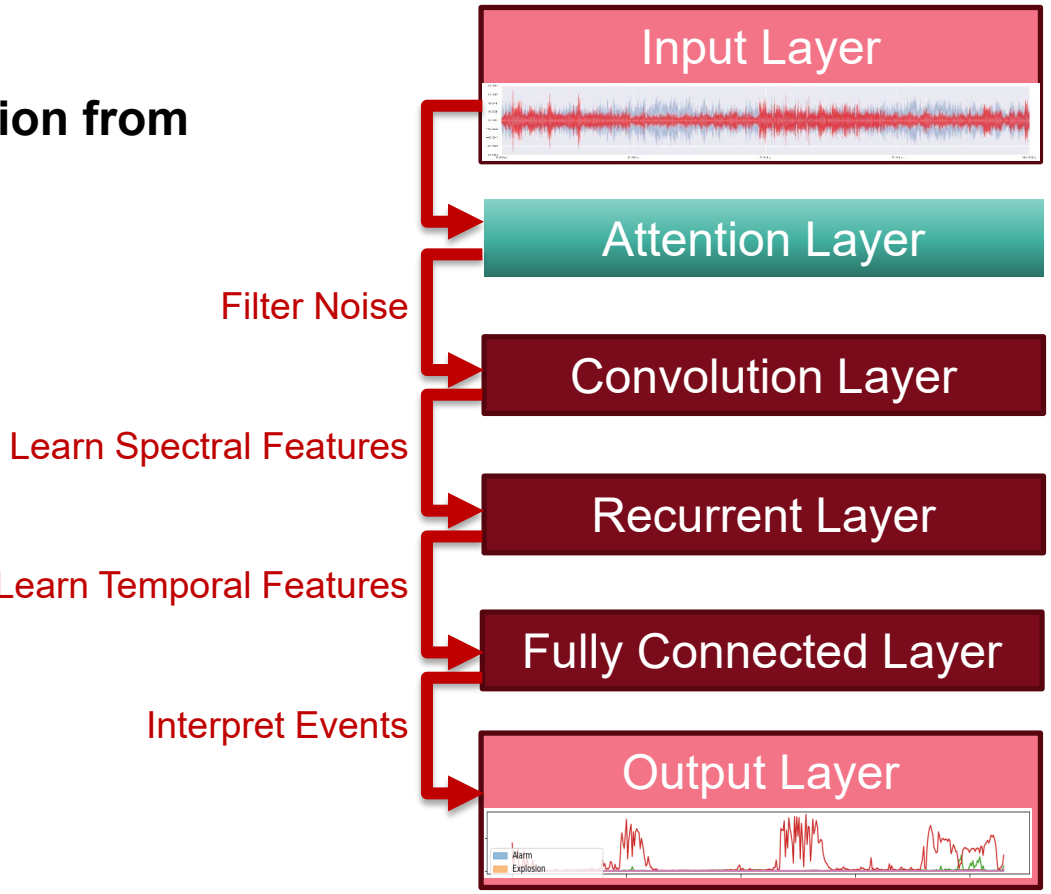

AUDIO EVENT DETECTION

1. Input representation
   - Common: Mel-Spectrograms
2. **Convolutional Neural Network Block (CNN)**
   - **Learn audio embeddings**

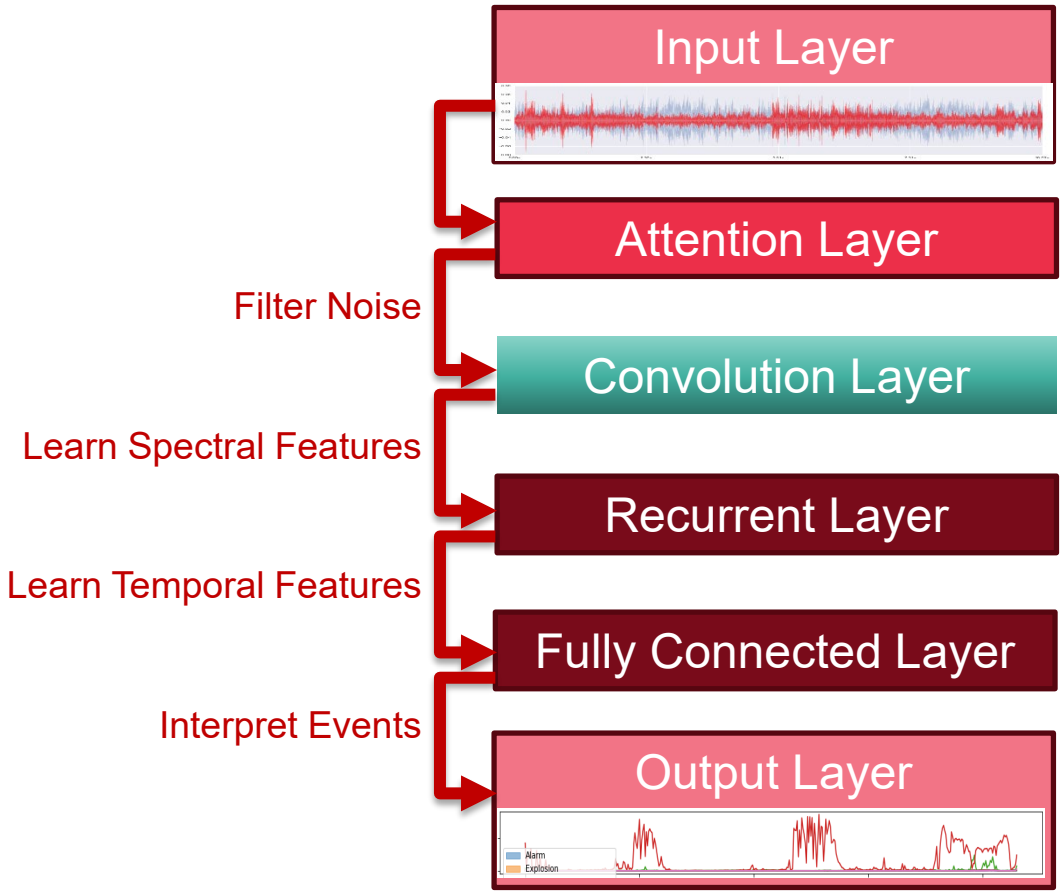

AUDIO EVENT DETECTION

1. Input representation
   - Common: Mel-Spectrograms
2. Convolutional Neural Network Block (CNN)
   - Learn audio embeddings
3. **Recurrent Neural Network Block (RNN)**
   - **Learn Temporal dependencies of embeddings**

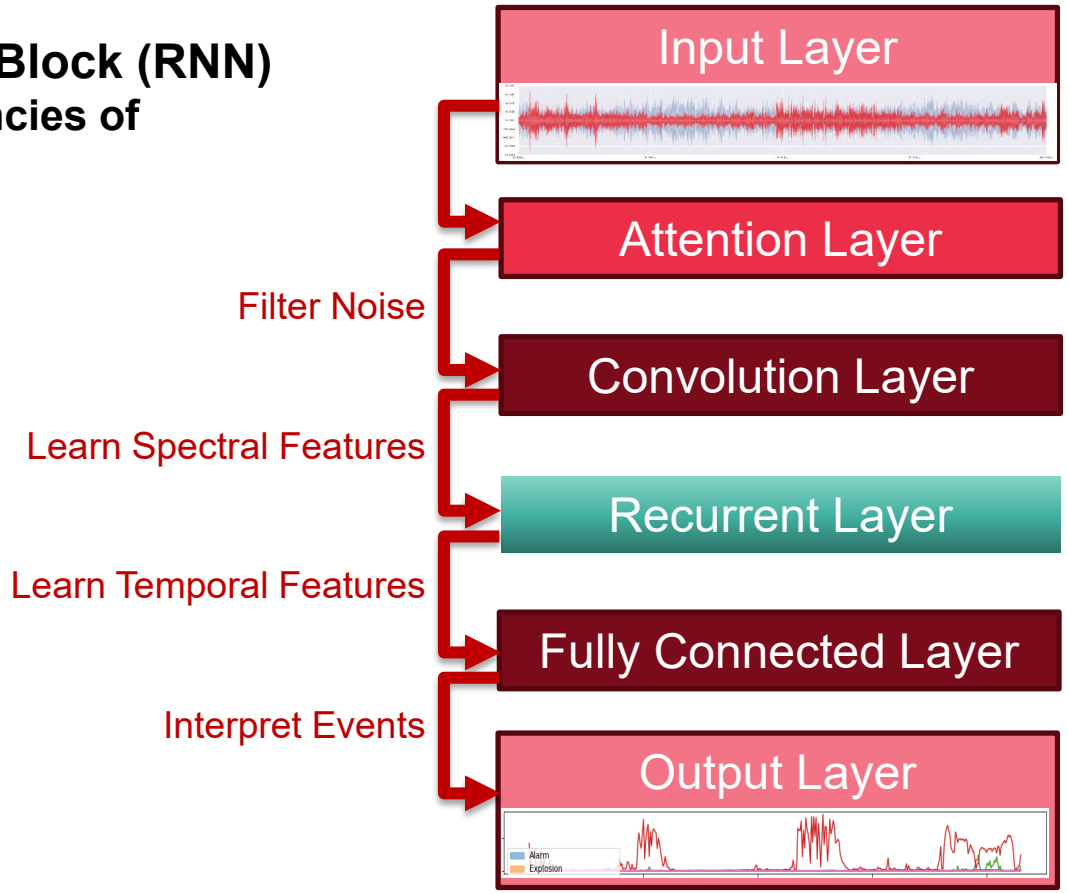

AUDIO EVENT DETECTION

1. Input representation
   - Common: Mel-Spectrograms
2. Convolutional Neural Network Block (CNN)
   - Learn audio embeddings
3. Recurrent Neural Network Block (RNN)
   - Learn Temporal dependencies of embeddings
4. **Array of Fully Connected Layers**
   - **One Layer per temporal dimension (Time-Distributed)**
   - **Dimensionality of Layer = Number of classes**

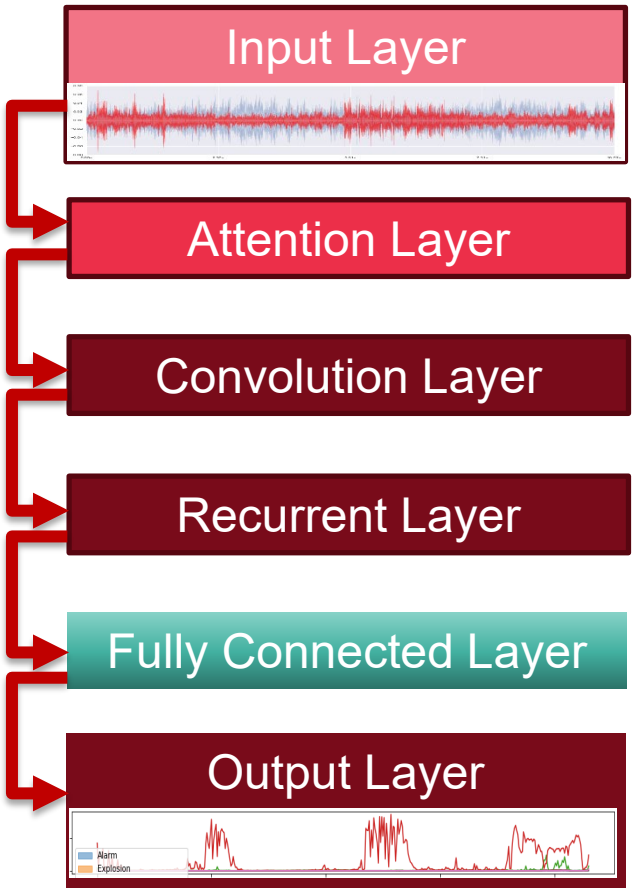

AUDIO EVENT DETECTION

1. Input representation
   - Common: Mel-Spectrograms
2. Convolutional Neural Network Block (CNN)
   - Learn audio embeddings
3. Recurrent Neural Network Block (RNN)
   - Learn Temporal dependencies of embeddings
4. Array of Fully Connected Layers
   - One Layer per temporal dimension (Time-Distributed)
   - Dimensionality of Layer = Number of classes
5. **Outputs**
   - **Strong Labels – Training & Inference**
     - Output of Time-Distributed Fully Connected Layers
   - **Weak Labels - Training**
     - Output Layer aggregation (e.g. avg, max)
     - Multi label prediction

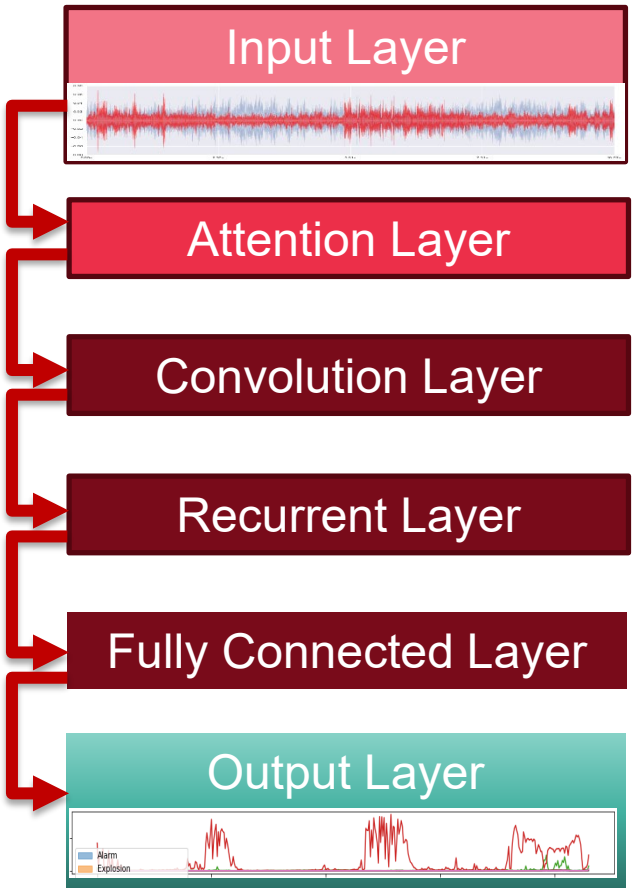

Google Audio Set

- 2M Videos
- 632 audio events
- annotaded according acoustic categories
- Weakly labelled (10s)
- Currently largest source of data

Human sounds
- Human voice
- Whistling
- Respiratory sounds
- Human locomotion
- Digestive
- Hands
- Heart sounds, heartbeat
- Otoacoustic emission
- Human group actions

Source-ambiguous sounds
- Generic impact sounds
- Surface contact
- Deformable shell

Animal
- Domestic animals, pets
- Livestock, farm animals, working animals
- Wild animals

Sounds of things
- Vehicle
- Engine
- Domestic sounds, home sounds
- Bell
- Alarm
- Mechanisms
- Tools
- Explosion

Music
- Musical instrument
- Music genre
- Musical concepts
- Music role
- Music mood

Natural sounds
- Wind
- Thunderstorm
- Water
- Fire

Channel, environment and background
- Acoustic environment
- Noise

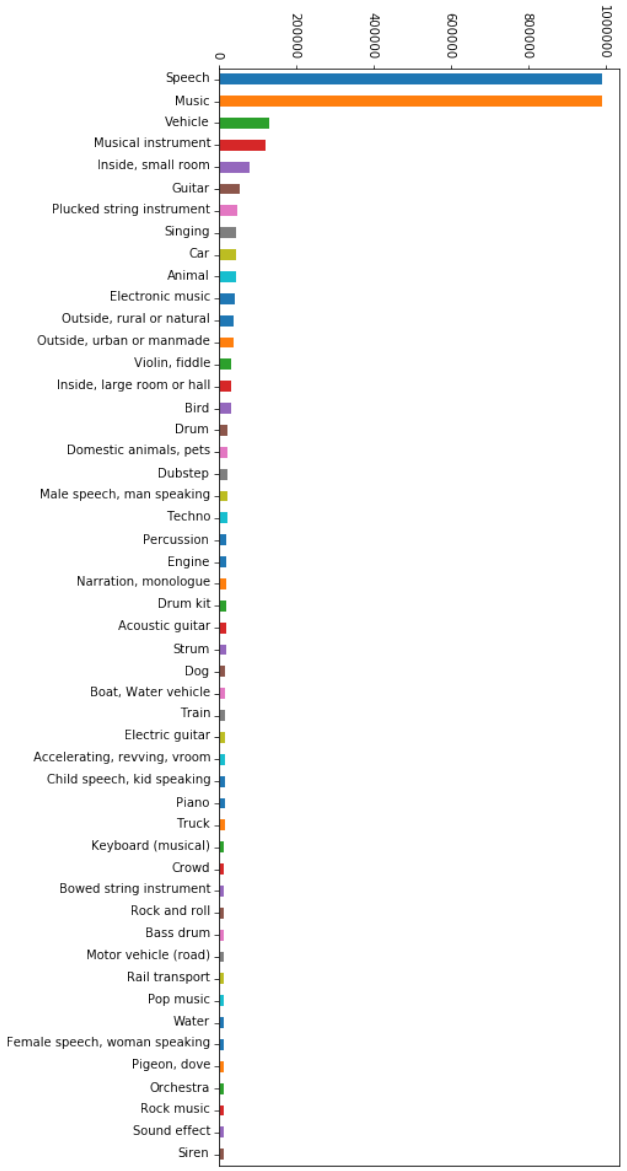

RECURRENT CONVOLUTIONAL NEURAL NETWORKS

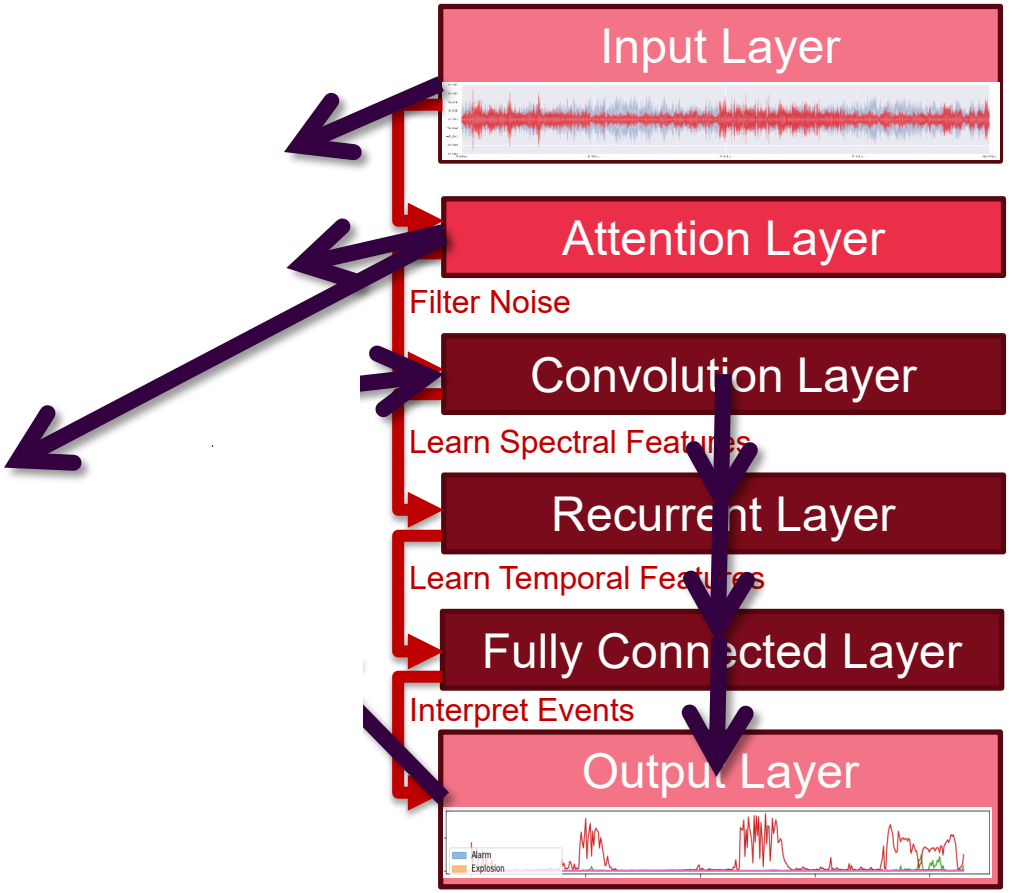

High signal to noise ratio
Due to attention layer

Smoothing Function
Temporal Segmentation per Event Category

AUDIO EVENTS VISUALIZED

User Interface Prototyp

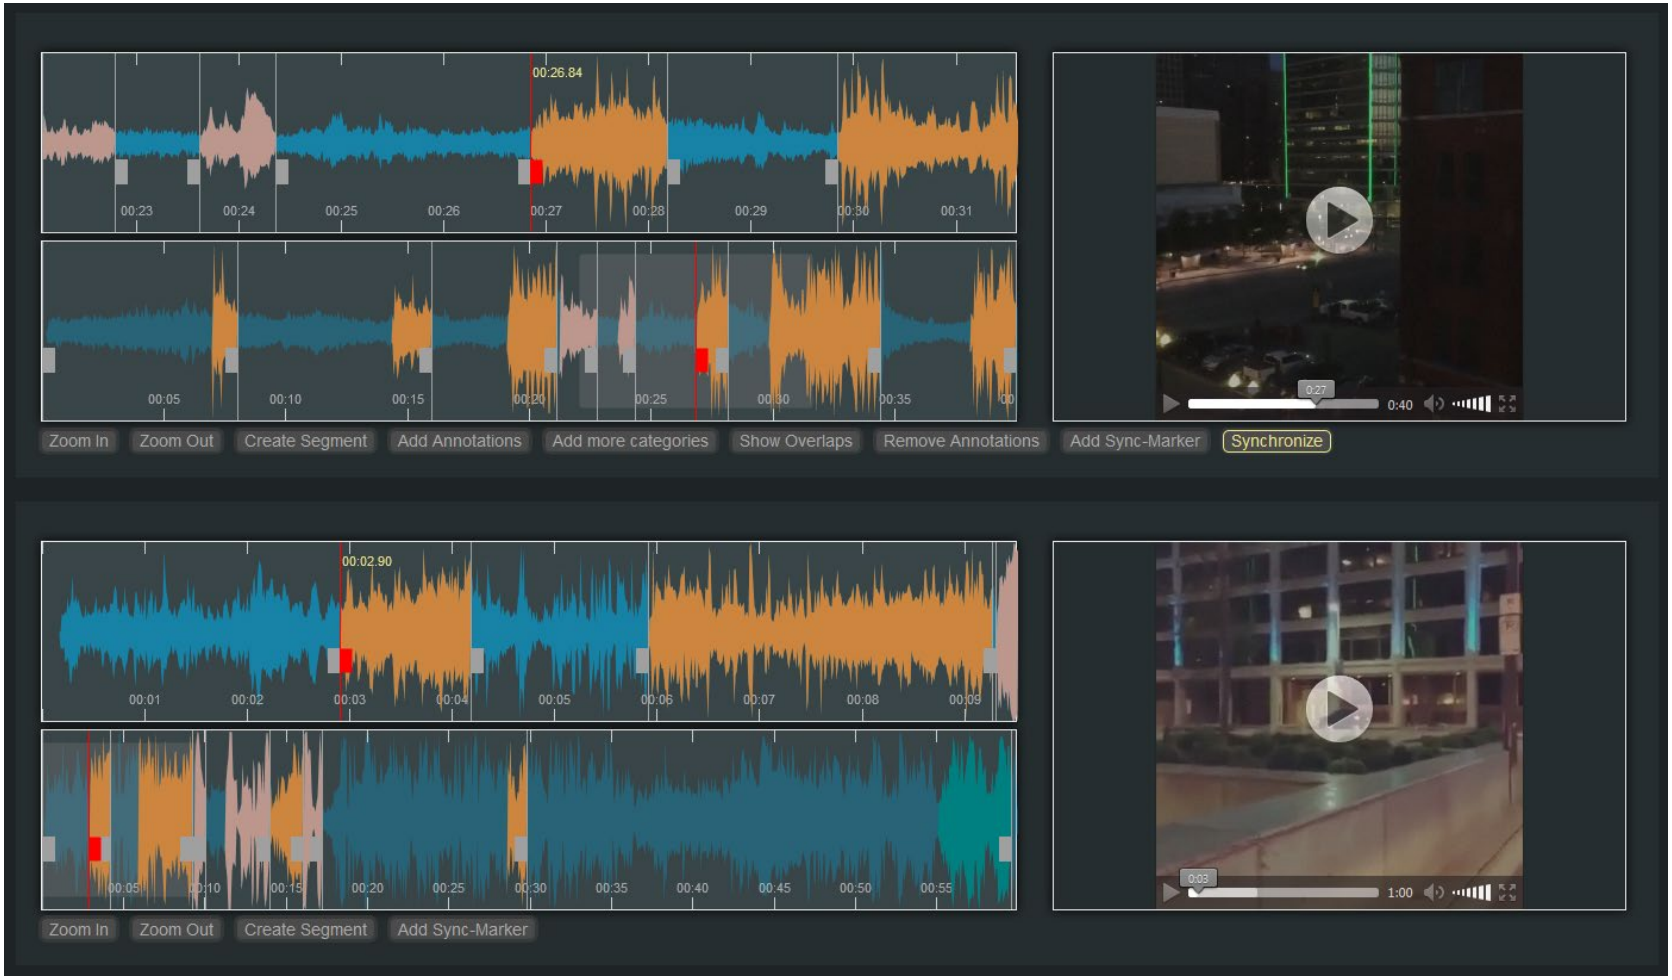

AUSTRIAN INSTITUTE OF TECHNOLOGY
TOMORROW TODAY

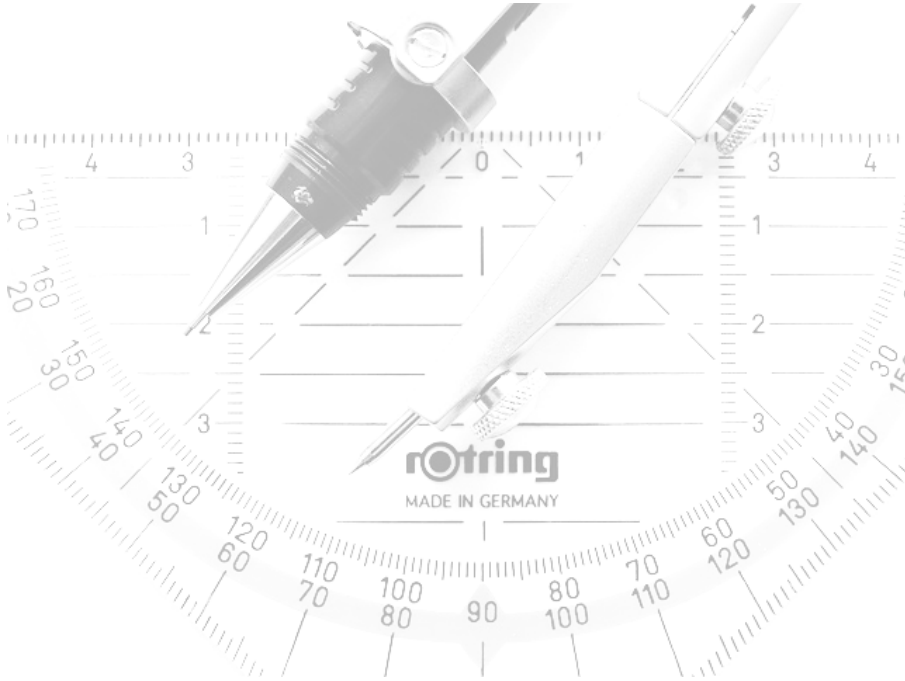

AUDIO SIMILARITY

Analytic Modules

AUDIO SIMILARITY SEARCH

- **Task**
  - Searching for video-segments with similar audio-signature
  - Sub-Segment video-search
- **Use-Case**
  - Suspect could not be identified in one video
  - Select segment and search for others using audio-signature
  - Instant localization (videos close to audio source)
- **Solution**
  - Select range in video
  - Retrieve a list of similar sounding video segments
  - Sorted by simlarity

AUDIO SIMILARITY SEARCH

- **Audio features extracted for each 6s segment**
  - Rhythm Patterns (repetitiveness in audio)
  - Statistical Spectrum Descriptors
- Nearest Neighbor search using late fusion in a normalized feature space

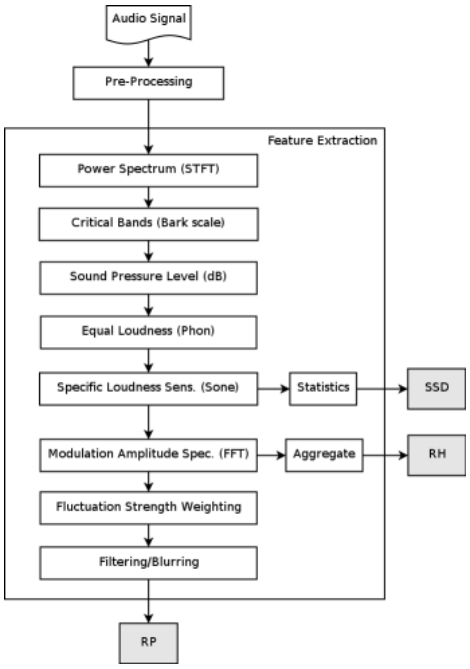

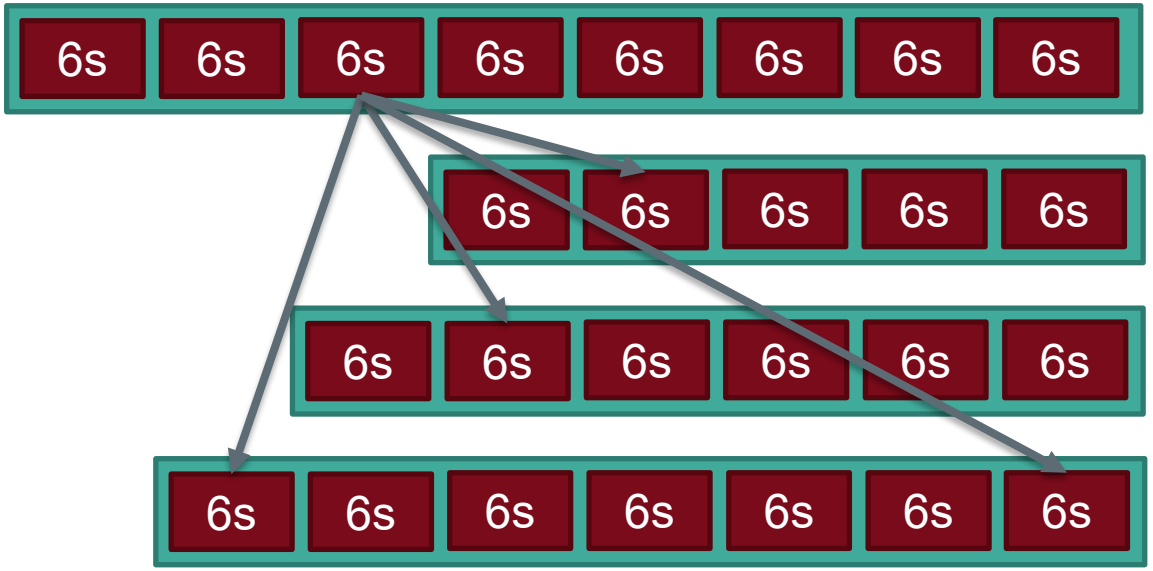

AUDIO SIMILARITY – INSTANT LOCALIZATION (WITHOUT GPS)

AUSTRIAN INSTITUTE OF TECHNOLOGY

Dallas Protest Shooting (2016)

AUSTRIAN INSTITUTE OF TECHNOLOGY
TOMORROW TODAY

AUDIO-BASED VIDEO-SYNCHRONIZATION

Analytic Module

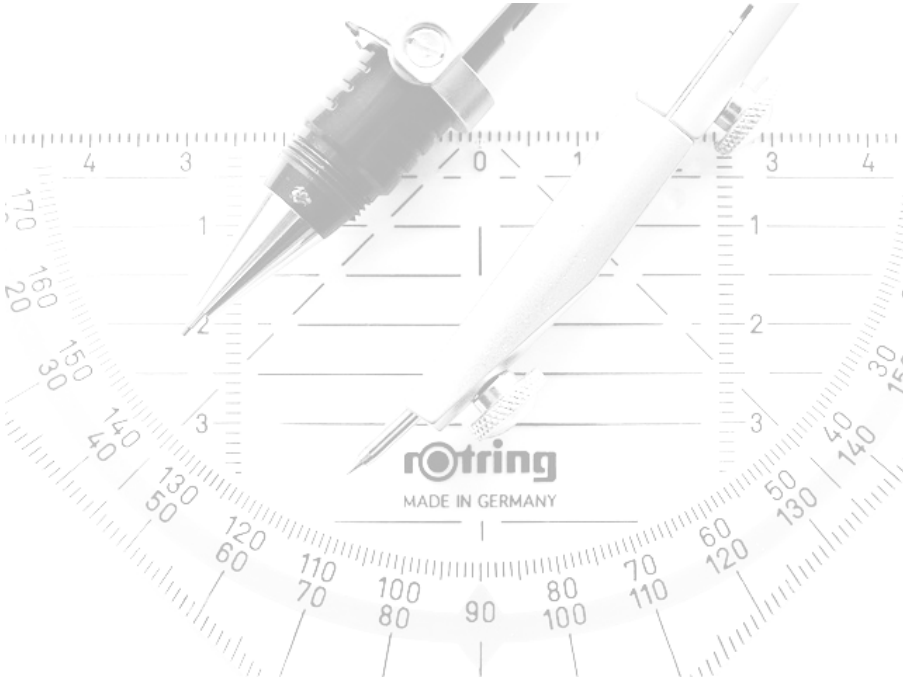

AUDIO-BASED VIDEO-SYNCHRONIZATION

- **Task**
  - Synchronize various video files with unreliable time metadata
  - Use audio-signature to relatively align video files

- **Technology**
  - Audio-fingerprints (chromaprint)
  - Noise invariant

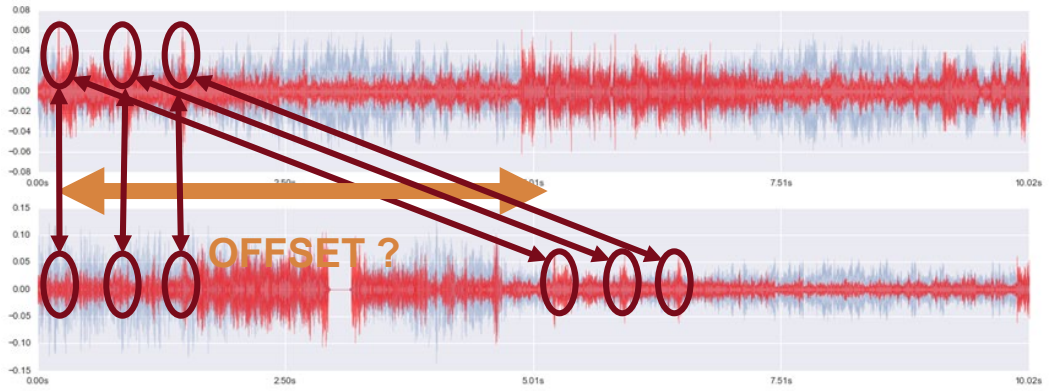

AUSTRIAN INSTITUTE OF TECHNOLOGY
TOMORROW TODAY

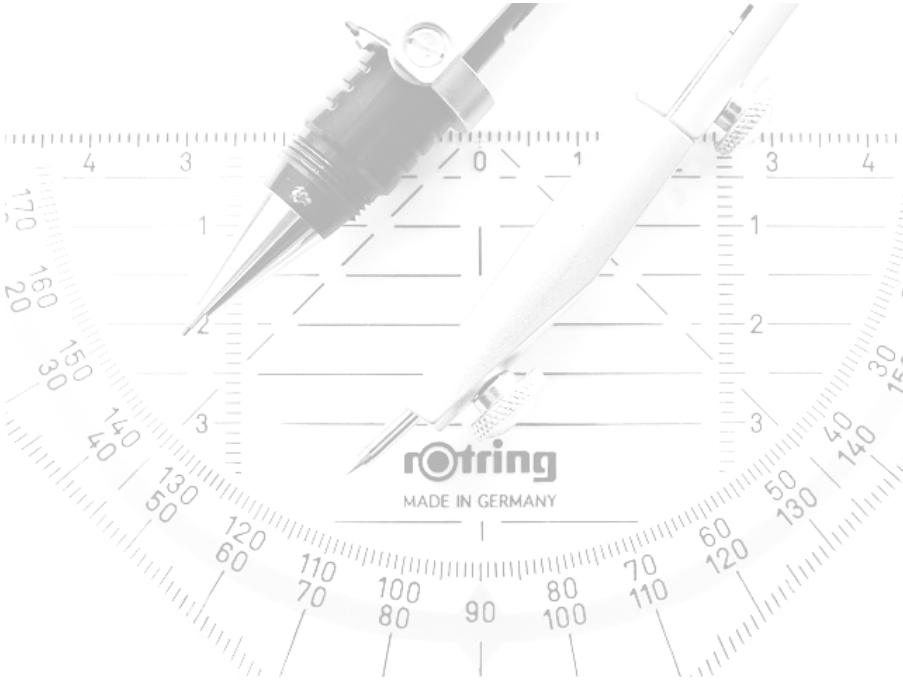

VISUAL ANALYTICS

Analytic Modules

VISUAL CONCEPT DETECTION

- YOLO
- License Plate detection
- Vehicle Color detection
- **Connected Vision** framework
  - Modular
  - Serviceoriented
  - Distributed
  - Scalable
- Rest Interface

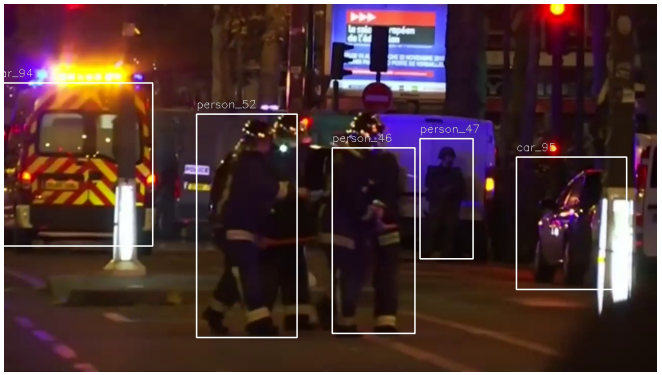

AUSTRIAN INSTITUTE OF TECHNOLOGY
TOMORROW TODAY

LARGE SCALE PLATFORM

Integration

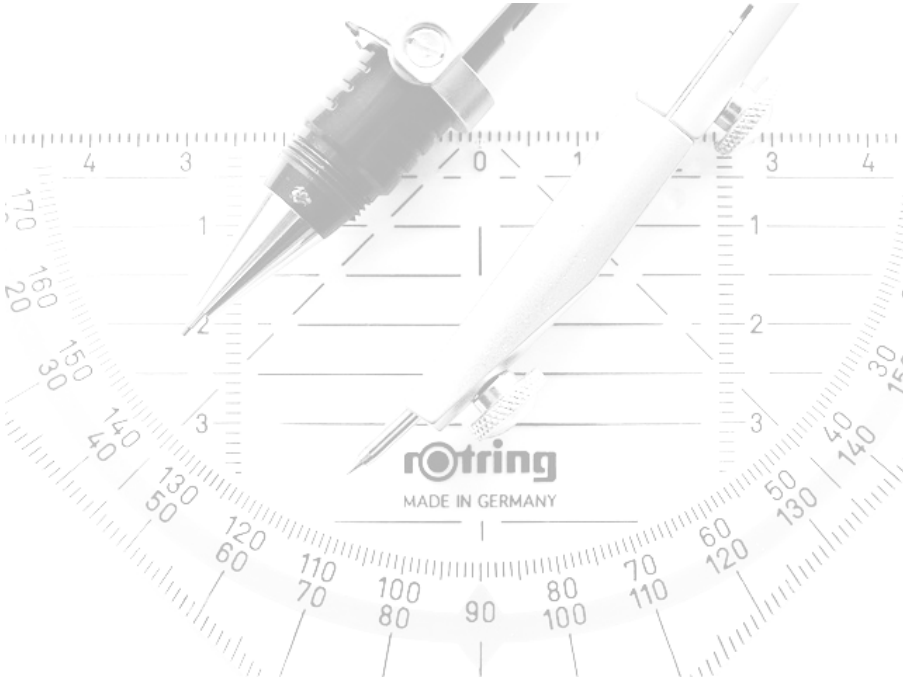

LARGE SCALE PLATFORM

Forensic Analytics in Massive Video Content

Visualization
Data Aggregation & Visualization
Interactive Processing

Large Scale Computing
Distributed Computing Clusters
Cloud Computing Big Data Architectures

Scalable Plattform
Apache Hadoop

GPU Plattform
NVIDIA, Cuda

Hive

Visual Object Detection
Suspects, Cars, Suitcases, Weapons
License plates, 3D Trajectories

Audio Event Detection
Gunshots, Explosions, Screams
Spoken Words, Transcription

Massive Video Data
e.g. after an terroristic attack
CCTV, mobile phones

GPU Scale Computing
Deep Learning, Deep Neural Networks
Artificial Intelligence Modules

CONCLUSIONS

- Audio Modules facilitate a rapid start into an investigation
  - Audio Event Detection => Filter
  - Audio Simlarity => Search
- Visual Modules facilitate a broad search for certain objects
  - Follow hints
- Hadoop is not the best choice for multimedia processing
- Integration of audio algorithms into pre-existing visual analytics systems holds pitfalls

THANK YOU

If you are interested
in the demo
ask me in the break!

AUSTRIAN INSTITUTE OF TECHNOLOGY

ALEXANDER SCHINDLER
Scientist
Information Management
Center for Digital Safety & Security

AIT Austrian Institute of Technology GmbH
Donau-City-Straße 1 | 1220 Wien
T +43 50550-2902 | M +43 664 8251454 | F +43 50550-2813
alexander.schindler@ait.ac.at | www.ait.ac.at